# The Bayesian Infinitesimal Jackknife for Variance: Supplemental information and proofs

Ryan Giordano (rgiordano@berkeley.edu) and
Tamara Broderick (tbroderick@mit.edu)

# A A simple closed-form example

In this section, we illustrate verification of our assumptions and the consistency of the IJ covariance with a simple closed form example, that of inferring a univariate normal mean with a misspecified variance. More elaborate examples can be found in Appendix H, including a Poisson model in the presence of overdispersed data and multilinear regression in the presence of heteroskedasticity.

For some fixed $\theta_{\text{True}} \in \mathbb{R}$ and $\sigma^2 \neq 1$, let us consider the following data generating process and model:

$$\begin{aligned} &\text{Data generating process:} && x_n \overset{iid}{\sim} \mathcal{N}(\cdot|\theta_{\text{True}}, \sigma^2) && \\ &\text{Model:} && x_n|\theta \overset{iid}{\sim} \mathcal{N}(\cdot|\theta, 1) && \theta \sim \mathcal{N}(\cdot|0, 1). \end{aligned}$$

We will take $\theta$ to be our quantity of interest, so $g(\theta) = \theta$. Since $\sigma^2 \neq 1$ the model is misspecified, and so we expect the scaled posterior variance $N \operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})}(\theta)$ to differ from $V^g$, even asymptotically.

For this simple model, the posterior is available in closed form:

$$\mathbb{P}(\theta|\mathcal{X}) = \mathcal{N}\left(\theta|\tilde{x}, (N+1)^{-1}\right) \quad \text{where} \quad \tilde{x} := \frac{1}{N+1}\sum_{n=1}^{N} x_n.$$

From this, we can see that $\theta_\infty = \lim_{N\to\infty} \operatorname*{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta] = \theta_{\text{True}}$. It happens in this simple example that $\theta_\infty$ is "consistent" for the true mean $\operatorname*{\mathbb{E}}_{\mathbb{F}(x_*)}[x_*]$, though in general this need not be the case. For example, the model's trivial "estimate" of $\operatorname*{Var}_{\mathbb{F}(x_*)}(x_*)$ is 1, which is never equal to the true value of $\sigma^2$.

One can readily verify that Assumptions 1 and 2 are satisfied for this model and this particular value of $\theta_\infty$. Arguably the most challenging condition is the final part of Assumption 1 Item 6, which one can show via strong concavity of the log likelihood, together with convergence of $\frac{1}{N}\sum_{n=1}^{N} \ell_{(1)}(x_n|\theta_\infty) \xrightarrow[N\to\infty]{prob} 0$ by the ordinary law of large numbers. Using special structure of a particular problem (like strong concavity) to establish a condition like Assumption 1 Item 6 is common in the literature — see, for example, textbook proofs of the consistency of M-estimators on unbounded domains (e.g. Theorem 9.11 of Keener (2010) or Chapter 6.2 of (Huber, 2011)).

We begin by establishing Corollary 1 and identifying our target, $V^g$. By the ordinary

CLT applied to our closed-form expression for $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta]$, we have that

$$\sqrt{N}\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta] - \theta_\infty\right) \xrightarrow[N\to\infty]{dist} \mathcal{N}\left(\cdot|0,\sigma^2\right),$$

from which we see that $V^g = \sigma^2$. It happens, in this case, that the variance of the limiting distribution is also the limit of the variance, since

$$N\mathop{\mathrm{Var}}_{\mathbb{F}(\mathcal{X})}\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta]\right) = \frac{N^2}{(N+1)^2}\sigma^2 \to \sigma^2 = V^g.$$

The asymptotic equivalence between the variance of the limit and the limit of the variance occurs in this case because the rescaled squared errors $N\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta] - \theta_\infty\right)^2$ are uniformly integrable in $N$ (Theorem 10.3.6 of Dudley (2018a)). We emphasize, however, that our theorems do not assume uniform integrability, and we aim to estimate the variance of the limiting normal distribution, not the limiting value of the variance, as discussed in the text surrounding Eq. 2.

To complete Corollary 1, we need to show that the expression for $V^g$ is correct. In this case, $g_{(1)}(\theta) = 1$, and the remaining quantities are (up to constants not depending on $\theta$):

$$\begin{aligned}
&\log\pi(\theta) = -\frac{1}{2}\theta^2 + \text{Constant} \quad \text{and} \quad \hat{\mathscr{L}}(\theta) = -\frac{1}{2N}\sum_{n=1}^{N}\left(x_n^2 - 2x_n\theta + \theta^2\right) - \frac{1}{2N}\theta^2 + \text{Constant} \Rightarrow \\
&\quad \hat{\theta} = \tilde{x} \quad , \quad \ell_{(1)}(x_n|\theta) = x_n - \theta \quad \text{and} \quad \hat{\mathscr{L}}_{(2)}(\theta) = -\frac{N+1}{N} \Rightarrow \\
&\quad \hat{\Sigma} = \frac{1}{N}\sum_{n=1}^{N}(x_n - \tilde{x})^2 \quad , \quad \hat{\mathcal{I}} = \frac{N+1}{N} \quad \text{and} \quad \hat{V}^g = \hat{\mathcal{I}}^{-1}\hat{\Sigma}\hat{\mathcal{I}}^{-1} = \frac{N^2}{(N+1)^2}\hat{\Sigma}. \qquad \text{(S1)}
\end{aligned}$$

It is readily verified that $\hat{V}^g \xrightarrow[N\to\infty]{prob} \sigma^2 = V^g$, so we have verified the conclusion of Corollary 1.

Finally, we manually verify Theorem 2. The influence score is given by:

$$\begin{aligned}
\psi_n &= N\mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X})}\left(\theta, -\frac{1}{2}\left(x_n^2 - 2x_n\theta + \theta^2\right)\right) \\
&= -\frac{1}{2}N\mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X})}(\theta, 1)\, x_n^2 + N\mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X})}(\theta,\theta)\, x_n - \frac{1}{2}N\mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X})}\left(\theta,\theta^2\right) = \frac{N}{N+1}x_n + \text{Constant},
\end{aligned}$$

where the constant in the final line is the same for all $x_n$. Since the IJ covariance depends only on the difference $\psi_n - \bar{\psi}$, quantities that are the same for each $n$ can be neglected in the IJ covariance computation. Therefore the IJ covariance estimate is given by

$$V^{\mathrm{IJ}} = \frac{1}{N}\sum_{n=1}^{N}\left(\psi_n - \bar{\psi}\right)^2 = \frac{N^2}{(N+1)^2}\frac{1}{N}\sum_{n=1}^{N}(x_n - \bar{x})^2 \xrightarrow[N\to\infty]{prob} \sigma^2 = V^g, \quad \text{where} \quad \bar{x} := \frac{1}{N}\sum_{n=1}^{N}x_n.$$

Since $\hat{\Sigma} = \frac{1}{N}\sum_{n=1}^{N}(x_n - \bar{x})^2 + (\bar{x} - \tilde{x})^2$ and $\bar{x} - \tilde{x} = \bar{x}/(N+1) = O_p(N^{-1})$, $V^{\mathrm{IJ}}$ differs from

the $\hat{V}^g$ given in Eq. S1 only by a term of order $O_p(N^{-2})$. So, like $\hat{V}^g$, $V^{\mathrm{IJ}}$ is a consistent estimator of the asymptotic variance $V^g$.

# B Notation key

For ease of reference, we provide here a table of key quantities defined and used through the course of our proofs.

| Symbol | Description | Definition |
|---|---|---|
| $\psi_n$ | Influence score of the $n$–th observation | Eq. 3 |
| $\mathcal{S}$ | A set of sets of data indices | Definition S1 |
| $\mathcal{X}_s$ | The observations corresponding to $s \in \mathcal{S}$ | Definition S1 |
| $\phi(\theta, \mathcal{X}_s)$ | A function of interest depending on $\theta$ and a set of datapoints | Definition S2 |
| $B_\delta$ | A $\delta$-ball around $\theta_\infty$ | Assumption 1 |
| $\mathscr{T}_{(k)}\left(\phi, \hat{\theta}, \theta\right)$ | The $k$–th order Taylor series residual term for $\phi$ | Definition S3 |
| $\hat{\mathcal{G}}(\tau)$ | Gaussian density with covariance $\hat{\mathcal{I}}^{-1}$ | Definition S8 |
| $\mathcal{G}_\lambda^S(\tau)$ | Isotropic Gaussian density with precision $\lambda$ truncated to set $S$ | Definition S8 |
| $\hat{Z}$ | Normalizing constant for $\hat{\mathcal{G}}(\tau)$ | Definition S8 |
| $\lambda_{\mathcal{I}}, \hat{\lambda}_{\mathcal{I}}$ | Minimum eigenvalues of $\mathcal{I}, \hat{\mathcal{I}}$ | Definition S4 |
| $\tau$ | Centered and scaled parameter $\sqrt{N}(\theta - \hat{\theta})$ | Eq. S11 |
| $R_1, R_2, R_3$ | Inner, middle, and outer regions of integration | Definition S7 |
| $\delta_1$ | Scale of boundary between $R_1$ and $R_2$ | Definition S7 |
| $\delta_2$ | Scale of boundary between $R_2$ and $R_3$ | Definition S7 |
| $I(\phi)$ | The integration functional of $\phi$ | Eq. S15 |
| $I_k(\phi)$ | Integration functional restricted to region $R_k$ | Eq. S16 |
| $\mathfrak{R}$ | The event that the problem is "regular" | Lemma S4 |
| $\epsilon_U$ | Arbitrarily small accuracy parameter for the regular event | Lemma S4 |
| $\epsilon_L$ | Log likelihood lower bound in regular event | Lemma S4 |
| $\mathscr{R}^\phi(\mathcal{X}_s)$ | The BCLT residual for $\mathop{\mathbb{E}}_{\mathbb{P}(\theta\|\mathcal{X})}[\phi(\theta, \mathcal{X}_s)]$ | Definition S9 |
| $\mathscr{R}^k(\mathcal{X}_s)$ | The $k$–th component of the BCLT residual | Definition S9 |

Table S1: Notation key

# C Setup and Preliminaries

## C.1 Derivative arrays

Throughout our proofs, we will be forming and manipulating higher-order Taylor expansions of functions of $\theta$. The arrays of derivatives in these expressions will always be summed against a direction, typically $\tau = \sqrt{N}(\theta - \hat{\theta})$. In order to avoid cumbersome summation notation for all such terms, we will implicitly sum derivative arrays against $\tau$ or $\theta$. For

example, for a function $\psi(\theta, x_n) \in \mathbb{R}^{D_\psi}$, we will write

$$\psi_{(3)}(\hat{\theta}, x_n)\tau^3 = \sum_{a=1}^{D}\sum_{b=1}^{D}\sum_{c=1}^{D} \left.\frac{\partial^3 \psi(\theta, x_n)}{\partial\theta_a \partial\theta_b \partial\theta_c}\right|_{\hat{\theta}} \tau_a\tau_b\tau_c.$$

Occasionally we will need to sum several derivative arrays against $\tau$, particularly in Section D.7. Since the directions $\tau$ will always be the same, we will combine the $\tau$ in such expressions with no ambiguity. For example, for the term involving $\hat{\mathscr{L}}_{(3)}(\hat{\theta})$ in the statement of Theorem 1, we will write

$$\psi_{(1)}(\hat{\theta}, x)\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^4 = \sum_{a=1}^{D}\sum_{b=1}^{D}\sum_{c=1}^{D}\sum_{d=1}^{D} \psi_{(1:a)}(\hat{\theta}, x)\hat{\mathscr{L}}_{(3:b,c,d)}(\hat{\theta})\tau_a\tau_b\tau_c\tau_d.$$

When necessary, higher derivatives will be indexed in the parentheses. For example, $\psi_{(2:i,j)}(\theta, x_n)$ will denote the length–$D_\psi$ vector

$$\left.\frac{\partial^2 \psi(\theta, x_n)}{\partial\theta_i \partial\theta_j}\right|_{\theta}.$$

We denote by $\mathbb{I}(A)$ the indicator function taking value 1 when the event $A$ occurs and 0 otherwise.

We will use order notation to describe the behavior of quantities as $N \to \infty$. For example, we will use $z = O(\phi(N))$ to mean there exists an $N^*$ and an $A > 0$ such that $N > N^* \Rightarrow z < A\phi(N)$. In our notation, $A$ and $N^*$ in this definition will never depend on the data, $\mathcal{X}$, though they may depend on unobserved population quantities. Dependence on the data $\mathcal{X}$ in order notation will always be represented with $O_p$ notation. For example, $z = O_p(\phi(N))$ will mean that there exists an $A > 0$ such that for any $\epsilon > 0$, there is a corresponding $N^*$ and such that

$$N > N^* \quad \Rightarrow \quad \mathop{\mathbb{E}}_{\mathbb{F}(\mathcal{X})}\left[\mathbb{I}\left(z < A\phi(N)\right)\right] > 1 - \epsilon.$$

Here, $A$ does not depend on $\epsilon$. The meaning is that the event $z = O(\phi(N))$ occurs with arbitrarily high probability as $N$ goes to infinity. The notation $\tilde{O}$ and $\tilde{O}_p$ will have the same meaning as $O$ and $O_p$ notation respectively, but with factors of $\log N$ possibly ignored.

## C.2 General index sets

To succinctly state our theorems for double sums, it will be convenient to introduce a more general index notation.

**Definition S1.** Let $\mathcal{S}$ denote a set of sets of data indices, where each member $s \in \mathcal{S}$ has the same cardinality, and for $s \in \mathcal{S}$, let $\mathcal{X}_s$ denote the set of corresponding datapoints. We write $\phi(\mathcal{X}_s)$ for a function that depends on $|s|$ datapoints. For example, if $\mathcal{S} = \{1, \ldots, N\} =: [N]$, then $\mathcal{X}_s = x_s$, and $\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}} \phi(\mathcal{X}_s) = \frac{1}{N}\sum_{n=1}^{N} \phi(x_n)$. Similarly, if $\mathcal{S} = [N] \times [N]$, then $\mathcal{X}_s = (x_{s_1}, x_{s_2})$, and $\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}} \phi(\mathcal{X}_s) = \frac{1}{N^2}\sum_{n=1}^{N}\sum_{m=1}^{N} \phi(x_n, x_m)$. ◀

We now state a generalization of Definition 2 that applies to generic index sets as defined

in Definition S1.

**Definition S2.** For an index set $\mathcal{S}$, a data-dependent function $\phi(\theta, \mathcal{X}_s)$ is $K$-th order BCLT-okay if

1. $\theta \mapsto \phi(\theta, \mathcal{X}_s)$ is $K$ times continuously differentiable $\mathbb{F}$–almost surely.
2. $\frac{1}{|\mathcal{S}|} \sum_{s \in \mathcal{S}} \mathop{\mathbb{E}}_{\pi(\theta)} \left[ \|\phi(\theta, \mathcal{X}_s)\|_2^2 \right] = O_p(1)$.
3. For some $\delta > 0$, and for each $k = 0, \ldots, K$, $\sup_{\theta \in B_\delta} \frac{1}{|\mathcal{S}|} \sum_{s \in \mathcal{S}} \left\| \phi_{(k)}(\theta, \mathcal{X}_s) \right\|_2^2 = O_p(1)$.

◀

Note that a function is BCLT-okay according to Definition 2 precisely if it is BCLT-okay according to Definition S2 with index set $\mathcal{S} = [N]$. Although $\mathcal{S} = [N]$ and $\mathcal{S} = [N] \times [N]$ will be the only examples of $\mathcal{S}$ we will use in the present paper, we believe it is clearer to keep the notation general.

In particular, Theorem 1 is stated in terms of $\mathcal{S} = [N]$, but we will prove that it holds for generic index sets of the form Definition S2. Formally, we will prove Theorem 1 under the following Assumption S1, noting that Assumption S1 is assumed by Theorem 1 for $\mathcal{S} = [N]$.

**Assumption S1.** *Assume that $\phi(\theta, \mathcal{X}_s)$ is third-order BCLT-okay according to Definition S2, and that Assumption 1 holds.*

## C.3 Taylor series expansions and their asymptotics

The proofs below repeatedly rely on two key ideas: Taylor series expansions, and uniform asymptotic control on the terms that arise. In this section, we introduce these ideas and prove two key lemmas. The first, Lemma S1, gives the form of a Taylor series expansion. The second, Lemma S2, controls the asymptotic behavior of the log likelihood and its derivatives in a neighborhood of $\theta_\infty$.

### C.3.1 Taylor series expansions

Our analysis is based strongly on Taylor series expansions. We will use the following two minor variants of standard analysis tools. The first is a multivariate version of the integral form of the Taylor series residual for a vector-valued function.

**Definition S3.** For a $k$-times differentiable function $\phi(\theta)$, a fixed $\hat{\theta}$, and $\theta$, define the residual functional

$$\mathscr{T}_{(k)}\left(\phi, \hat{\theta}, \theta\right) := \frac{1}{(k-1)!} \int_0^1 (1-t)^{k-1} \phi_{(k)}(\hat{\theta} + t(\theta - \hat{\theta})) dt.$$

◀

The residual $\mathscr{T}_{(k)}\left(\phi, \hat{\theta}, \theta\right)$ is an array of the same dimension as the $k$-th order derivative, $\phi_{(k)}(\theta)$. Observe that, for any norm $\|\cdot\|$,

$$\left\| \mathscr{T}_{(k)}\left(\phi, \hat{\theta}, \theta\right) \right\| \leq \frac{1}{(k-1)!} \sup_{\tilde{\theta}: \left\|\tilde{\theta} - \hat{\theta}\right\|_\infty \leq \left\|\theta - \hat{\theta}\right\|_\infty} \left\| \phi_{(k)}(\tilde{\theta}) \right\|.$$

**Lemma S1.** *A $K$-times continuously differentiable function $\phi(\theta)$ has Taylor series residual*

$$\phi(\theta) = \sum_{k=0}^{K-1} \frac{1}{k!} \phi_{(k)}(\hat{\theta})(\theta - \hat{\theta})^k + \mathscr{T}_{(K)}\left(\phi, \hat{\theta}, \theta\right)(\theta - \hat{\theta})^K.$$

*Proof.* Let $\phi_d(\theta)$ denote the $d$-th component of the vector $\phi(\theta)$ (contrast this notation with the $d$-th order derivative $\phi_{(d)}(\theta)$). Consider the scalar to scalar map $t \mapsto \phi_d(\hat{\theta} + t(\theta - \hat{\theta}))$. Taking the integral remainder form of the $K-1$-th order Taylor series expansion of this scalar map (Dudley, 2018b, Appendix B.2) and stacking the result into a vector gives the desired result. □

### C.3.2 Log likelihood asymptotics

To use Lemma S1 to study $\hat{\mathscr{L}}(\theta)$, we will need to expand Definition 1 with the following quantities.

**Definition S4.**

$$\hat{\mathscr{L}}_{(k)}(\theta) := \frac{1}{N}\sum_{n=1}^{N} \ell_{(k)}(x_n|\theta) + \frac{1}{N}\left(\log \pi\right)_{(k)}(\theta) \quad \mathscr{L}_{(k)}(\theta) := \int \ell_{(k)}(x_*|\theta) d\mathbb{F}(x_*)$$

$$\lambda_{\mathcal{I}} := \text{The minimum eigenvalue of } \mathcal{I} \qquad \hat{\lambda}_{\mathcal{I}} := \text{The minimum eigenvalue of } \hat{\mathcal{I}}$$

Note that, following our convention, $\hat{\mathscr{L}}_{(k)}(\theta)$ is the $k$–th partial derivative of $\hat{\mathscr{L}}(\theta)$ with respect to $\theta$.

The term $\mathscr{L}_{(k)}(\theta)$ is a potential abuse of notation, since it is only the partial derivative of $\mathscr{L}(\theta)$ under the assumption that we can interchange integration and differentiation. Though we will in fact be assuming sufficient conditions for this interchange (Assumption 1 Item 5), our proofs will simply use the above definition of $\mathscr{L}_{(k)}(\theta)$ directly. ◀

Using Lemma S1, we can then form a Taylor series expansion around $\hat{\theta}$, an expression that we will analyze carefully in Appendix D:

$$\begin{aligned}\hat{\mathscr{L}}(\theta) = \hat{\mathscr{L}}(\hat{\theta}) + \hat{\mathscr{L}}_{(1)}(\hat{\theta})(\theta - \hat{\theta}) + \frac{1}{2}\hat{\mathscr{L}}_{(2)}(\hat{\theta})(\theta - \hat{\theta})^2 + \\ \frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})(\theta - \hat{\theta})^3 + \mathscr{T}_{(4)}\left(\hat{\mathscr{L}}, \hat{\theta}, \theta\right)(\theta - \hat{\theta})^4. \end{aligned} \tag{S2}$$

We now turn to the task of controlling the kinds of terms that arise in Eq. S2. One of our key tools in the proof of Lemma S2 will be *uniform laws of large numbers* (ULLNs), a concept which we make concrete in Definition S5. Typically, ULLNs require that the sample average $\frac{1}{N}\sum_{n=1}^{N}\phi(x_n, \theta)$ converge to $\mathop{\mathbb{E}}_{\mathbb{F}(x_*)}[\phi(x_n, \theta)]$ uniformly over $\theta$ in some compact set.

**Definition S5.** Recall the definition of $B_\delta$ from Assumption 1. We say a "uniform law of large numbers", or "ULLN", holds for a function $\phi(\theta, x_*)$ if, for some $\delta > 0$,

$$\sup_{\theta \in B_\delta} \left\| \frac{1}{N}\sum_{n=1}^{N} \phi(\theta, x_n) - \mathop{\mathbb{E}}_{\mathbb{F}(x_*)}[\phi(\theta, x_*)] \right\|_2 \xrightarrow[N\to\infty]{prob} 0,$$

where $\mathop{\mathbb{E}}_{\mathbb{F}(x_*)}[\phi(\theta, x_*)]$ is finite. ◀

In general, a ULLN requires that the set of functions $\{\theta \mapsto \phi(\theta, \mathcal{X}) : \|\theta - \theta_\infty\|_2 < \delta\}$ be a "Glivenko-Cantelli class" with finite covering number (Chapter 19 of van der Vaart (2000), van der Vaart and Wellner (2013)). A simple condition for a uniform law of large numbers to hold is that $\phi(\theta, x_*)$ be dominated by an absolutely integrable function of $x_*$ alone in a compact neighborhood of $\theta_\infty$ (Example 19.7 of van der Vaart (2000)).

We now state a key lemma controlling the asymptotic behavior of the log likelihood and its derivatives.

**Lemma S2.** *Suppose that there exists $\delta_{LLN} > 0$, $M(x_*)$, and $K \geq 1$ such that*

$$\sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K)}(x_*|\theta)\right\|_2 \leq M(x_*) \quad \textit{with} \quad \mathop{\mathbb{E}}_{\mathbb{F}(x_*)}\left[M(x_*)^2\right] < \infty. \tag{S3}$$

*Additionally assume that $\pi(\theta)$ is non-zero and $K-1$ times continuously differentiable in $B_{\delta_{LLN}}$. Assume also that*

$$\mathop{\mathbb{E}}_{\mathbb{F}(x_*)}\left[\left\|\ell_{(k)}(x_*|\theta_\infty)\right\|_2^2\right] < \infty \textit{ for } k = 0, \ldots, K-1. \tag{S4}$$

*Then the following results hold. First, for $k \in \{0, \ldots, K\}$,*

$$\sup_{\theta \in B_{\delta_{LLN}}} \frac{1}{N}\sum_{n=1}^{N} \left\|\ell_{(k)}(x_n|\theta)\right\|_2^2 = O_p(1). \tag{S5}$$

*Then, for $k \in \{0, \ldots, K-1\}$,*

$$\sup_{\theta \in B_{\delta_{LLN}}} \left\| \frac{1}{N}\sum_{n=1}^{N}\left(\ell_{(k)}(x_n|\theta) - \mathop{\mathbb{E}}_{\mathbb{F}(x_*)}\left[\ell_{(k)}(x_*|\theta)\right]\right)\right\|_2 \xrightarrow[N\to\infty]{prob} 0 \tag{S6}$$

$$\sup_{\theta \in B_{\delta_{LLN}}} \left\|\hat{\mathscr{L}}_{(k)}(\theta) - \mathscr{L}_{(k)}(\theta)\right\|_2 \xrightarrow[N\to\infty]{prob} 0. \tag{S7}$$

*Proof.* First, note that

$$\sup_{\theta \in B_{\delta_{LLN}}} \frac{1}{N}\sum_{n=1}^{N} \left\|\ell_{(K)}(x_n|\theta)\right\|_2^2 \leq \frac{1}{N}\sum_{n=1}^{N} \sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K)}(x_n|\theta)\right\|_2^2 \quad \text{and}$$

$$\mathop{\mathbb{E}}_{\mathbb{F}(x_*)}\left[\sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K)}(x_*|\theta)\right\|_2^2\right] \leq \mathop{\mathbb{E}}_{\mathbb{F}(x_*)}\left[M(x_*)^2\right] < \infty,$$

so Eq. S5 holds for $k = K$ by the strong law of large numbers.

We now prove the remaining relations inductively. For a scalar–valued, differentiable

function $f(\theta)$, and any two $\theta, \theta'$ in $B_{\delta_{LLN}}$, we can write

$$\begin{aligned}
|f(\theta) - f(\theta')| &= \left| \int_0^1 \left. \frac{\partial f(\theta' + t(\theta - \theta'))}{\partial t} \right|_t dt \right| \\
&= \left| \int_0^1 f_{(1)}(\theta' + t(\theta - \theta'))(\theta - \theta') dt \right| \\
&\leq \|\theta - \theta'\|_2 \sup_{\tilde{\theta} \in B_{\delta_{LLN}}} \left\| f_{(1)}(\tilde{\theta}) \right\|_2 .
\end{aligned}$$

Applying the previous formula componentwise to $\ell_{(K-1)}(x_*|\theta)$, and loosening the right hand side to contain all of $\left\|\ell_{(K)}(x_*|\tilde{\theta})\right\|_2$ rather than only a particular component, gives

$$\begin{aligned}
\left\|\ell_{(K-1)}(x_*|\theta) - \ell_{(K-1)}(x_*|\theta')\right\|_\infty &\leq \|\theta - \theta'\|_2 \sup_{\tilde{\theta} \in B_{\delta_{LLN}}} \left\|\ell_{(K)}(x_*|\tilde{\theta})\right\|_2 \\
&\leq \|\theta - \theta'\|_2 M(x_*). \qquad \text{(S8)}
\end{aligned}$$

It follows that each component of $\ell_{(K-1)}(x_*|\theta)$ is Glivenko-Cantelli, and so Eq. S6 holds for $k = K - 1$, using also Eq. S3 and Eq. S4 for square integrability (see van der Vaart, 2000, Example 19.7, as well as the discussion in Definition S5 above). Further, evaluating Eq. S8 at $\theta' = \theta_\infty$ and taking the supremum gives

$$\begin{aligned}
\sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K-1)}(x_*|\theta)\right\|_2^2 &\leq D_\theta^{K-1} \sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K-1)}(x_*|\theta)\right\|_\infty^2 \\
&\leq 2D_\theta^{K-1} \sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K-1)}(x_*|\theta) - \ell_{(K-1)}(x_*|\theta_\infty)\right\|_\infty^2 + 2D_\theta^{K-1} \left\|\ell_{(K-1)}(x_*|\theta_\infty)\right\|_\infty^2 \\
&\leq 2D_\theta^{K-1} \sup_{\theta \in B_{\delta_{LLN}}} \|\theta - \theta_\infty\|_2^2 M(x_*)^2 + 2D_\theta^{K-1} \left\|\ell_{(K-1)}(x_*|\theta_\infty)\right\|_2^2 \\
&\leq 2\delta_{LLN}^2 D_\theta^{K-1} M(x_*)^2 + 2D_\theta^{K-1} \left\|\ell_{(K-1)}(x_*|\theta_\infty)\right\|_2^2 .
\end{aligned}$$

It follows from Eq. S3 and Eq. S4 with $k = K - 1$ that

$$\mathop{\mathbb{E}}_{\mathbb{F}(x_*)} \left[ \sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K-1)}(x_*|\theta)\right\|_2^2 \right] < \infty,$$

and so Eq. S5 holds for $k = K - 1$ by applying the LLN to the envelope function $\sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K-1)}(x_*|\theta)\right\|_2^2$.

Since Eq. S3 holds for $k = K-1$ with the envelope $M_{K-1}(x_*) := \left(\sup_{\theta \in B_{\delta_{LLN}}} \left\|\ell_{(K-1)}(x_*|\theta)\right\|_2^2\right)^{1/2}$ in place of $M(x_*)$, the remaining results follow by proceeding inductively on $K - 1$, $K - 2$, and so on.

Finally, by continuity of $\pi(\theta)$ (Assumption 1 Item 2) and compactness of $B_{\delta_{LLN}}$, we have that, as $N \to \infty$,

$$\frac{1}{N} \sup_{\theta \in B_{\delta_{LLN}}} \left\| (\log \pi)_{(k)} (\theta) \right\|_2^2 \to 0.$$

Then Eq. S7 follows from the preceding display, Cauchy-Schwarz, and Eq. S6. □

## C.4 Consistency results

Finally, we prove known results on the frequentist behavior of key quantities using our own assumptions and notation.

**Lemma S3.** *Under Assumption 1,* $\hat{\theta} \xrightarrow[N\to\infty]{prob} \theta_\infty$*, and* $\left\|\hat{\mathcal{I}}^{-1}\right\|_{op} < 2\lambda_{\mathcal{I}}^{-1}$*, with probability approaching one.*

*Proof.* First, by strict optimality of the log likelihood (Assumption 1 Item 6) and the boundedness of the prior density (Assumption 1 Item 2), we must have $\hat{\theta} \in B_\delta$ with probability approaching one as $N \to \infty$, since

$$
\begin{aligned}
&\inf_{\theta\in\Omega_\theta\setminus B_\delta} \left(\hat{\mathscr{L}}(\theta_\infty) - \hat{\mathscr{L}}(\theta)\right) \\
&\quad= \inf_{\theta\in\Omega_\theta\setminus B_\delta} \left(\frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta_\infty) - \frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta) + \frac{1}{N}\log\pi(\theta_\infty) - \frac{1}{N}\log\pi(\theta)\right) \\
&\quad\geq \inf_{\theta\in\Omega_\theta\setminus B_\delta} \left(\frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta_\infty) - \frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta)\right) + \\
&\qquad \frac{1}{N}\left(\log\pi(\theta_\infty) - \sup_{\theta\in\Omega_\theta} \mathbb{I}\left(\pi(\theta)\geq 1\right)\log\pi(\theta)\right) \\
&\quad\geq \varepsilon + \frac{1}{N}\left(\log\pi(\theta_\infty) - \sup_{\theta\in\Omega_\theta} \mathbb{I}\left(\pi(\theta)\geq 1\right)\log\pi(\theta)\right) && \text{(Assumption 1 Item 6)} \\
&\quad\geq \varepsilon + O\left(N^{-1}\right). && \text{(Assumption 1 Item 2)}
\end{aligned}
$$

By definition of $\hat{\theta}$, $\hat{\mathscr{L}}(\theta_\infty) - \hat{\mathscr{L}}(\hat{\theta}) \leq 0$, and so $\hat{\theta} \notin \Omega_\theta \setminus B_\delta$ for sufficiently large $N$. It therefore suffices to consider the behavior of $\hat{\theta} \in B_\delta$.

Next, fix any $\delta' \in (0, \delta)$. We will show that, conditioned on $\hat{\theta} \in B_\delta$, with probability approaching one, we must also have $\hat{\theta} \in B_{\delta'}$. Define

$$3\eta := \inf_{\theta\in B_\delta\setminus B_{\delta'}} \left(\mathscr{L}(\theta_\infty) - \mathscr{L}(\theta)\right) > 0.$$

We know that $\eta > 0$ since the closure of $B_\delta \setminus B_{\delta'}$ is compact, $\mathscr{L}(\theta)$ is continuous on $B_\delta$, $\theta_\infty$ is not in the closure of $B_\delta \setminus B_{\delta'}$, and that $\theta_\infty$ is the unique maximum of $\mathscr{L}(\theta)$ (Assumption 1 Items 4 and 6). Note that $\eta$ depends only on $\delta$, $\delta'$, and population parameters, and not on the data. By a ULLN (Lemma S2 with $k = 0$), we can choose $N$ large enough so that, with arbitrarily high probability,

$$\sup_{\theta\in B_\delta} \left|\hat{\mathscr{L}}(\theta) - \mathscr{L}(\theta)\right| < \eta. \tag{S9}$$

Then, on the event that Eq. S9 holds,

$$\begin{aligned}
&\inf_{\theta \in B_\delta \setminus B_{\delta'}} \left( \hat{\mathscr{L}}(\theta_\infty) - \hat{\mathscr{L}}(\theta) \right) \\
&\geq -\left| \hat{\mathscr{L}}(\theta_\infty) - \mathscr{L}(\theta_\infty) \right| + \inf_{\theta \in B_\delta \setminus B_{\delta'}} \left( \mathscr{L}(\theta_\infty) - \mathscr{L}(\theta) \right) - \sup_{\theta \in B_\delta \setminus B_{\delta'}} \left| \mathscr{L}(\theta) - \hat{\mathscr{L}}(\theta) \right| \\
&\geq \inf_{\theta \in B_\delta \setminus B_{\delta'}} \left( \mathscr{L}(\theta_\infty) - \mathscr{L}(\theta) \right) - 2 \sup_{\theta \in B_\delta} \left| \hat{\mathscr{L}}(\theta) - \mathscr{L}(\theta) \right| \\
&\geq 3\eta - 2\eta > 0.
\end{aligned}$$

However, as argued above, $\hat{\mathscr{L}}(\theta_\infty) - \hat{\mathscr{L}}(\hat{\theta}) \leq 0$, and so when Eq. S9 holds, $\hat{\theta} \notin B_\delta \setminus B_{\delta'}$. So it must be the case that $\hat{\theta} \in B_{\delta'}$, and since $\delta'$ was arbitrary, it must be the case that $\hat{\theta} \xrightarrow[N \to \infty]{prob} \theta_\infty$.

It remains to bound $\left\| \hat{\mathcal{I}}^{-1} \right\|_{op}$, where $\hat{\mathcal{I}} = -\hat{\mathscr{L}}_{(2)}(\hat{\theta})$. By Assumption 1 Item 6, $\mathcal{I}$ is positive definite with $\|\mathcal{I}^{-1}\|_{op} = \left\| \mathscr{L}_{(2)}(\theta_\infty)^{-1} \right\|_{op} = \lambda_{\mathcal{I}}^{-1}$. Since $\delta'$ above was arbitrary, we may choose $\delta'$ small enough that, by continuity of the operator norm,

$$\sup_{\theta \in B_{\delta'}} \left\| \mathscr{L}_{(2)}(\theta)^{-1} \right\|_{op} < \sqrt{2} \lambda_{\mathcal{I}}^{-1}.$$

Next, by the ULLN of Assumption 1 Item 5 and Lemma S2,

$$\sup_{\theta \in B_{\delta'}} \left\| \hat{\mathscr{L}}_{(2)}(\theta)^{-1} \right\|_{op} < 2\lambda_{\mathcal{I}}^{-1},$$

with probability approaching one. Since $\hat{\theta} \in B_{\delta'}$ with probability approaching one, it follows that $\left\| \hat{\mathcal{I}}^{-1} \right\|_{op} < 2\lambda_{\mathcal{I}}^{-1}$ with probability approaching one.

□

**Proof of Proposition 1.**

*Proof.* By Lemma S3, $\hat{\theta} \in B_\delta$ with probability approaching one. On this event, using Lemma S1 to Taylor expand $\hat{\mathscr{L}}_{(1)}$ around $\hat{\theta}$ and evaluate at $\theta_\infty$ gives

$$\hat{\mathscr{L}}_{(1)}(\theta_\infty) = \hat{\mathscr{L}}_{(1)}(\hat{\theta}) + \hat{\mathscr{T}}_{(1)} \left( \hat{\mathscr{L}}_{(1)}, \hat{\theta}, \theta_\infty \right) (\theta_\infty - \hat{\theta}). \tag{S10}$$

Now, let $\delta_N := \left\| \hat{\theta} - \theta_\infty \right\|_2$, so that, by Lemma S3, $\delta_N \xrightarrow[N \to \infty]{prob} 0$. Then, conditional on

$\delta_N < \delta$, observe that

$$\begin{aligned}
&\left\| \mathscr{T}_{(1)}\left(\hat{\mathscr{L}}_{(1)}, \hat{\theta}, \theta_\infty\right) - \mathscr{L}_{(2)}(\theta_\infty) \right\|_2 \\
&= \left\| \int_0^1 \left( \hat{\mathscr{L}}_{(2)}(t(\theta_\infty - \hat{\theta}) + \hat{\theta}) - \mathscr{L}_{(2)}(\theta_\infty) \right) dt \right\|_2 \\
&\leq \sup_{\theta \in B_{\delta_N}} \left\| \hat{\mathscr{L}}_{(2)}(\theta) - \mathscr{L}_{(2)}(\theta_\infty) \right\|_2 \\
&\leq \sup_{\theta \in B_{\delta}} \left\| \hat{\mathscr{L}}_{(2)}(\theta) - \mathscr{L}_{(2)}(\theta) \right\|_2 + \sup_{\theta \in B_{\delta_N}} \left\| \mathscr{L}_{(2)}(\theta) - \mathscr{L}_{(2)}(\theta_\infty) \right\|_2 \\
&\xrightarrow[N\to\infty]{prob} 0,
\end{aligned}$$

where the limit to zero follows from the ULLN via Assumption 1 Item 5 and Lemma S2 for the first term, and, for the second term, from $\delta_N \xrightarrow[N\to\infty]{prob} 0$ together with the continuity of the non-random function $\mathscr{L}_{(2)}$ at $\theta_\infty$ (Assumption 1 Item 4). It follows that

$$\mathscr{T}_{(1)}\left(\hat{\mathscr{L}}_{(1)}, \hat{\theta}, \theta_\infty\right)^{-1} \xrightarrow[N\to\infty]{prob} \mathscr{L}_{(2)}(\theta_\infty)^{-1} = -\mathcal{I}^{-1}.$$

Next, by Assumption 1 Item 5, $\mathop{\mathbb{E}}_{\mathbb{F}(x)}\left[\left\|\ell_{(1)}(x|\theta_\infty)\right\|_2^2\right]$ is finite, so $\mathop{\mathrm{Cov}}_{\mathbb{F}(x)}\left(\ell_{(1)}(x|\theta_\infty)\right) = \Sigma$ is finite. By smoothness of $\mathscr{L}$ (Assumption 1 Item 4) and the definition of $\theta_\infty$, $\mathop{\mathbb{E}}_{\mathbb{F}(x)}\left[\ell_{(1)}(x|\theta_\infty)\right] = 0$. The $x_n$ are IID, so, by a central limit theorem,

$$\sqrt{N}\hat{\mathscr{L}}_{(1)}(\theta_\infty) = \frac{1}{\sqrt{N}}\sum_{n=1}^{N} \ell_{(1)}(x_n|\theta_\infty) + \frac{1}{\sqrt{N}}\left(\log \pi\right)_{(1)}(\theta_\infty) \xrightarrow[N\to\infty]{dist} \mathcal{N}(\cdot|0, \Sigma),$$

where the prior term is $O(N^{-1/2})$ because the prior is continuously differentiable and non-zero at $\theta_\infty$ (Assumption 1 Item 2), so the limit follows from Slutsky's theorem. By Eq. S10 and $\hat{\mathscr{L}}_{(1)}(\hat{\theta}) = 0$, with probability approaching one,

$$\begin{aligned}
\sqrt{N}(\hat{\theta} - \theta_\infty) = &- \mathscr{T}_{(1)}\left(\hat{\mathscr{L}}_{(1)}, \hat{\theta}, \theta_\infty\right)^{-1} \sqrt{N}\hat{\mathscr{L}}_{(1)}(\theta_\infty) \\
&\xrightarrow[N\to\infty]{dist} \mathcal{N}(\cdot|0, \mathcal{I}^{-1}\Sigma\mathcal{I}^{-1}),
\end{aligned}$$

where the final line follows from Slutsky's theorem. Finally, $g(\theta)$ is continuously differentiable by Assumption 2, so the delta method (van der Vaart, 2000, Theorem 3.1) gives

$$\sqrt{N}\left(g(\hat{\theta}) - g(\theta_\infty)\right) \xrightarrow[N\to\infty]{dist} \mathcal{N}\left(\cdot|0, g_{(1)}(\theta_\infty)\mathcal{I}^{-1}\Sigma\mathcal{I}^{-1}g_{(1)}(\theta_\infty)^\intercal\right) = \mathcal{N}\left(\cdot|0, V^g\right).$$

□

# D Proof of Theorem 1

The proof of Theorem 1 will be broken up into several parts. Our first step is to derive an explicit series expansion for posterior expectations of the form $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\phi(\theta)]$, culminating in Lemma S11. In Lemma S14, we then consider the dependence of $\phi(\theta)$ on a subset of datapoints, writing $\phi(\theta, \mathcal{X}_s)$ (recalling Definition S1), and showing that the results of Lemma S11 can be summed over $\mathcal{S}$. The residual of the expansion in Lemma S11 is designed so that Lemma S14 follows easily.

To prove Lemma S11, we follow classical proofs of the BCLT, keeping explicit track of the residual. We first express the desired expectation as the ratio of integrals with respect to $\tau$. We then divide the domain of integration into three regions according to their distance from $\hat{\theta}$ and treat each separately. Finally, we combine the bounds on the integrals over different regions for our final result.

Throughout, we will use the change of variable

$$\theta := \hat{\theta} + N^{-1/2}\tau \quad \Leftrightarrow \quad \tau = \sqrt{N}(\theta - \hat{\theta}). \tag{S11}$$

We will switch the domain of integration between $\tau$ and $\theta$ without re-writing the arguments. For example, for some function $\phi(\theta)$, we will write the shorthand $\int \phi(\theta) d\tau$ for the more precise $\int \phi\left(N^{-1/2}\tau + \hat{\theta}\right) d\tau$. Similarly, we will overload set notation: if $A$ is a set of $\theta$, we will write $\tau \in A$ as a shorthand for $\tau \in \left\{\tau : \hat{\theta} + N^{-1/2}\tau \in A\right\}$. Although slight abuses of notation, these conventions will keep our expressions more compact, as well as serve as a reminder of the $\theta$ domain in which the integrand is smooth by assumption.

## D.1 Ratio of integrals

We need to consider the asymptotic behavior of $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\phi(\theta, \mathcal{X}_s)]$, which can be written as the ratio of the integrals

$$\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\phi(\theta, \mathcal{X}_s)] = \frac{\int \phi(\theta, \mathcal{X}_s) \exp\left(N\hat{\mathscr{L}}(\theta)\right) d\theta}{\int \exp\left(N\hat{\mathscr{L}}(\theta)\right) d\theta}.$$

Recall that $\hat{\theta} := \text{argmax}_\theta \hat{\mathscr{L}}(\theta)$ and define $\tilde{\theta}(t) := t(\theta - \hat{\theta}) + \hat{\theta}$. We will employ the following Taylor series expansion:

$$\begin{aligned}\hat{\mathscr{L}}(\theta) = \hat{\mathscr{L}}(\hat{\theta}) + \hat{\mathscr{L}}_{(1)}(\hat{\theta})(\theta - \hat{\theta}) + \frac{1}{2}\hat{\mathscr{L}}_{(2)}(\hat{\theta})(\theta - \hat{\theta})^2 + \\ \frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})(\theta - \hat{\theta})^3 + \mathscr{T}_{(4)}\left(\hat{\mathscr{L}}, \hat{\theta}, \theta\right)(\theta - \hat{\theta})^4,\end{aligned} \tag{S12}$$

recalling the definition of the integral form of the Taylor series residual $\mathscr{T}_{(4)}\left(\hat{\mathscr{L}}, \hat{\theta}, \theta\right)$ given in Definition S3.

Note that $\hat{\mathscr{L}}(\hat{\theta})$ is constant in $\theta$, and that $\hat{\mathscr{L}}_{(1)}(\hat{\theta}) = 0$. Plugging in and changing the

domain of integration to $\tau$, we see

$$\begin{aligned}\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\phi(\theta,\mathcal{X}_s)] &= \frac{\int \phi(\theta,\mathcal{X}_s)\exp\left(N\hat{\mathscr{L}}(\theta)-N\hat{\mathscr{L}}(\hat{\theta})\right)d\theta}{\int \exp\left(N\hat{\mathscr{L}}(\theta)-N\hat{\mathscr{L}}(\hat{\theta})\right)d\theta} \\ &= \frac{\int \phi(\theta,\mathcal{X}_s)\exp\left(N\hat{\mathscr{L}}(\theta)-N\hat{\mathscr{L}}(\hat{\theta})\right)N^{-D_\theta/2}d\tau}{\int \exp\left(N\hat{\mathscr{L}}(\theta)-N\hat{\mathscr{L}}(\hat{\theta})\right)N^{-D_\theta/2}d\tau}.\end{aligned}$$

Since the factor of $N^{-D_\theta/2}$ cancels in the preceding display, it will suffice to consider the following generic integral.

**Definition S6.** For a target function $\phi$, define the integral functional

$$I(\phi) := \int \phi(\theta,\mathcal{X}_s)\exp\left(N\hat{\mathscr{L}}(\theta)-N\hat{\mathscr{L}}(\hat{\theta})\right)d\tau \tag{S13}$$

$$= \int \phi(\theta,\mathcal{X}_s)\times \tag{S14}$$

$$\exp\left(\frac{1}{2}\hat{\mathscr{L}}_{(2)}(\hat{\theta})\tau^2 + N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 + N^{-1}\mathscr{T}_{(4)}\left(\hat{\mathscr{L}},\hat{\theta},\theta\right)\tau^4\right)d\tau. \tag{S15}$$

◀

Under Definition S6, $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\phi(\theta,\mathcal{X}_s)] = I(\phi)/I(1)$. We will now consider the behavior of $I(\phi)$ for a generic third-order BCLT-okay $\phi$, returning to consider the behavior of the ratio in Section D.8.

## D.2 Regions

For some $\delta_1$ and $\delta_2$ to be determined later, we will divide consideration of $I(\phi)$ into the following three disjoint regions.

**Definition S7.** For any $\delta_1 > 0$ and $\delta_2 > 0$ and $N$ sufficiently large, define the disjoint regions

$$\begin{aligned}R_1 &:= \left\{\theta : \left\|\theta-\hat{\theta}\right\|_2 < \delta_1\frac{\log N}{\sqrt{N}}\right\} \\ R_2 &:= \left\{\theta : \delta_1\frac{\log N}{\sqrt{N}} \le \left\|\theta-\hat{\theta}\right\|_2 < \delta_2\right\} \\ R_3 &:= \left\{\theta : \delta_2 \le \left\|\theta-\hat{\theta}\right\|_2\right\}.\end{aligned}$$

Recalling our slight abuse of set notation, the same regions in the domain of $\tau$ are,

$$\begin{aligned}\tau \in R_1 &\Rightarrow \tau \in \{\tau : \|\tau\|_2 < \delta_1 \log N\} \\ \tau \in R_2 &\Rightarrow \left\{\tau : \delta_1 \log N \le \|\tau\|_2 < \delta_2\sqrt{N}\right\} \\ \tau \in R_3 &\Rightarrow \left\{\tau : \delta_2\sqrt{N} \le \|\tau\|_2\right\}.\end{aligned}$$

Similarly, we will use the regions as domains of integration for $\tau$ as well as $\theta$. Correspondingly, we can define

$$I_1(\phi) := \int_{R_1} \phi(\theta, \mathcal{X}_s) \exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) d\tau$$
$$I_2(\phi) := \int_{R_2} \phi(\theta, \mathcal{X}_s) \exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) d\tau$$
$$I_3(\phi) := \int_{R_3} \phi(\theta, \mathcal{X}_s) \exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) d\tau. \tag{S16}$$

◀

Since the boundaries of the regions coincide, $I(\phi) = I_1(\phi) + I_2(\phi) + I_3(\phi)$ whenever the regions $R_1$, $R_2$, and $R_3$ are disjoint, which will occur for $N$ sufficiently large (see Lemma S4 below for more discussion of this point). Note that the union $R_1 \bigcup R_2$ is a $\delta_2$ ball around $\hat{\theta}$, and that $R_1$ shrinks to $\hat{\theta}$ as $N \to \infty$.

Throughout the proof, we will be choosing $\delta_1$ and $\delta_2$ to satisfy certain conditions. Our assumptions will always take the form of assuming that $\delta_1$ is *large enough*, and that $\delta_2$ is *small enough*. For any particular $\delta_1$ and $\delta_2$, we then will require that $N$ is large enough that the regions are disjoint.

## D.3 Nearly certain events

Our entire analysis will be conditioned on the occurrence of several events, each of which occurs with probability approaching one as $N$ goes to infinity. Outside of these events we provide no guarantees. As a consequence, the residuals of our approximation can be used to show convergence in probability or distribution (via Slutsky's theorem), but not convergence in expectation nor almost sure convergence.

**Lemma S4.** *For any $\epsilon_U > 0$, let $\mathfrak{R}$ denote the event that all of the following occur for a particular $N$, $\delta_1$, $\delta_2$. ("R" is for "regular".)*

1. *The regions $R_1$, $R_2$, and $R_3$ are disjoint, and $I(\phi) = I_1(\phi) + I_2(\phi) + I_3(\phi)$.*
2. *$\left\|\hat{\theta} - \theta_\infty\right\|_2 < \epsilon_U$.*
3. *$|\pi(\hat{\theta}) - \pi(\theta_\infty)| < \epsilon_U$.*
4. *For each $k \in \{0, \ldots, 4\}$, $\sup_{\theta \in R_1 \bigcup R_2} \left\|\hat{\mathscr{L}}_{(k)}(\theta) - \mathscr{L}_{(k)}(\theta)\right\|_2 \leq \epsilon_U$.*
5. *$\hat{\lambda}_{\mathcal{I}} > \lambda_{\mathcal{I}}/2$ and $\left\|\hat{\mathcal{I}}\right\|_{op} < 2\left\|\mathcal{I}\right\|_{op}$.*
6. *There exists $\epsilon_L > 0$ such that $\inf_{\theta \in R_3}\left(\sum_{n=1}^N \ell(x_n|\hat{\theta}) - \sum_{n=1}^N \ell(x_n|\theta)\right) > N\epsilon_L$.*

*Let Assumption 1 hold with $\delta = \delta_{LLN}$. For any $\delta_1 > 0$, any $\pi_0 \in (0,1)$, and any $\delta_2 > 0$ and $\epsilon_U > 0$ satisfying $\delta_{LLN} > \epsilon_U + \delta_2$ and $\delta_2/2 \geq \epsilon_U$, there exists an $N^*$ such $N > N^*$ implies that event $\mathfrak{R}$ occurs with $\mathbb{F}$-probability at least $1 - \pi_0$.*

*Proof.* Event 1 follows from the fact that, for any positive $\delta_1$ and $\delta_2$, eventually $\delta_2\sqrt{N} > \delta_1 \log N$.

Event 2 follows from Lemma S3.

Event 3 follows from the fact that the prior density is continuous by Assumption 1 Item 2 and, by Lemma S3, we can take $\hat{\theta}$ to be arbitrarily close to $\theta_\infty$.

For Event 4, observe that $\left\|\hat{\theta} - \theta_\infty\right\|_2 < \epsilon_U$ implies that

$$\begin{aligned}
\theta \in & R_1 \bigcup R_2 \Rightarrow \\
\left\|\theta - \hat{\theta}\right\|_2 \leq & \delta_2 \Rightarrow \\
\|\theta - \theta_\infty\|_2 \leq \left\|\theta - \hat{\theta}\right\|_2 + \left\|\hat{\theta} - \theta_\infty\right\|_2 & \leq \delta_2 + \epsilon_U \leq \delta_{LLN}.
\end{aligned}$$

Therefore, for each $k \in \{0, \ldots, 4\}$,

$$\sup_{\theta \in R_1 \bigcup R_2} \left\|\hat{\mathscr{L}}_{(k)}(\theta) - \mathscr{L}_{(k)}(\theta)\right\|_2 \leq \sup_{\theta \in B_{\delta_{LLN}}} \left\|\hat{\mathscr{L}}_{(k)}(\theta) - \mathscr{L}_{(k)}(\theta)\right\|_2 \xrightarrow[N\to\infty]{prob} 0.$$

For $k \leq 3$, the limit is Eq. S7 of Lemma S2 with $K = 4$. For $k = 4$, $\theta \mapsto \ell_{(4)}(x_*|\theta)$ is continuous on the compact set $B_{\delta_{LLN}}$ with integrable envelope $M(x_*)$ (Assumption 1 Items 4 and 5), so it satisfies a ULLN (van der Vaart, 2000, Example 19.8), and $\frac{1}{N}\sup_{\theta \in B_{\delta_{LLN}}} \left\|(\log \pi)_{(4)}(\theta)\right\|_2 \to 0$ by Assumption 1 Item 2. Since there are finitely many $k$, all of these bounds hold simultaneously with probability approaching one.

Event 5 follows from Lemma S3, together with the fact that $\hat{\mathcal{I}} = -\hat{\mathscr{L}}_{(2)}(\hat{\theta}) \xrightarrow[N\to\infty]{prob} \mathcal{I}$ by $\hat{\theta} \xrightarrow[N\to\infty]{prob} \theta_\infty$, the ULLN of Lemma S2, the continuity of $\mathscr{L}_{(2)}$, and the continuity of matrix eigenvalues.

For Event 6, we have to convert Assumption 1 Item 6 into a statement about $\hat{\theta}$ and $R_3$. First,

$$\begin{aligned}
\inf_{\theta \in R_3} & \left(\frac{1}{N}\sum_{n=1}^N \ell(x_n|\hat{\theta}) - \frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta)\right) = \\
& \frac{1}{N}\sum_{n=1}^N \ell(x_n|\hat{\theta}) - \frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta_\infty) + \inf_{\theta \in R_3}\left(\frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta_\infty) - \frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta)\right). \quad \text{(S17)}
\end{aligned}$$

We begin by controlling the second term in Eq. S17. Since $\left\|\hat{\theta} - \theta_\infty\right\|_2 < \epsilon_U \leq \delta_2/2$ by Event 2 of the present lemma,

$$\begin{aligned}
\theta \in & R_3 \Rightarrow \\
\left\|\theta - \theta_\infty + \theta_\infty - \hat{\theta}\right\|_2 \geq & \delta_2 \Rightarrow \\
\|\theta - \theta_\infty\|_2 + \left\|\theta_\infty - \hat{\theta}\right\|_2 \geq & \delta_2 \Rightarrow \\
\|\theta - \theta_\infty\|_2 > & \delta_2 - \frac{\delta_2}{2} = \frac{\delta_2}{2},
\end{aligned}$$

It follows that the set $R_3$ is contained in $\Omega_\theta \setminus B_{\delta_2/2}$, and

$$\inf_{\theta \in R_3} \left( \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta_\infty) - \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta) \right) \geq \inf_{\theta \in \Omega_\theta \setminus B_{\delta_2/2}} \left( \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta_\infty) - \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta) \right).$$

Now, applying Assumption 1 Item 6, we have

$$\inf_{\theta \in \Omega_\theta \setminus B_{\delta_{LLN}}} \left( \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta_\infty) - \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta) \right) > \varepsilon,$$

for some $\varepsilon > 0$.

It remains to control the annulus

$$\begin{aligned}
&\inf_{\theta \in B_{\delta_{LLN}} \setminus B_{\delta_2/2}} \left( \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta_\infty) - \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta) \right) \\
&= \inf_{\theta \in B_{\delta_{LLN}} \setminus B_{\delta_2/2}} \left( \hat{\mathscr{L}}(\theta_\infty) - \hat{\mathscr{L}}(\theta) - \frac{1}{N} \log \pi(\theta_\infty) + \frac{1}{N} \log \pi(\theta) \right) \\
&\geq \inf_{\theta \in B_{\delta_{LLN}} \setminus B_{\delta_2/2}} \left( \hat{\mathscr{L}}(\theta_\infty) - \hat{\mathscr{L}}(\theta) \right) - \frac{2}{N} \sup_{\theta \in B_{\delta_{LLN}}} |\log \pi(\theta)|.
\end{aligned}$$

Define $\eta$ to satisfy

$$4\eta := \inf_{\theta \in B_{\delta_{LLN}} \setminus B_{\delta_2/2}} (\mathscr{L}(\theta_\infty) - \mathscr{L}(\theta)) > 0$$

Since $\theta_\infty$ is a unique maximizer of $\mathscr{L}$, $\mathscr{L}(\theta)$ is continuous, and the closure of $B_{\delta_{LLN}} \setminus B_{\delta_2/2}$ is a compact set not containing $\theta_\infty$, we have that $\eta > 0$.

We now bound the remaining terms by $\eta$. Since the prior is continuous and nonzero on the compact set $B_{\delta_{LLN}}$, and because $B_{\delta_{LLN}}$ is fixed by $\delta_{LLN}$, $\sup_{\theta \in B_{\delta_{LLN}}} |\log \pi(\theta)|$ is a finite constant and we can take $N^*$ large enough that

$$\frac{2}{N} \sup_{\theta \in B_{\delta_{LLN}}} |\log \pi(\theta)| < \eta. \tag{S18}$$

By Lemma S2 (with $k = 0$), there exists an $N^*$ such that, with arbitrarily high probability,

$$\sup_{\theta \in B_{\delta_{LLN}}} \left| \hat{\mathscr{L}}(\theta) - \mathscr{L}(\theta) \right| \leq \eta.$$

Combining,

$$\begin{aligned}
&\inf_{\theta \in B_{\delta_{LLN}} \setminus B_{\delta_2/2}} \left( \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta_\infty) - \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta) \right) \\
&\geq \inf_{\theta \in B_{\delta_{LLN}} \setminus B_{\delta_2/2}} \left( \hat{\mathscr{L}}(\theta_\infty) - \mathscr{L}(\theta_\infty) + \mathscr{L}(\theta_\infty) - \mathscr{L}(\theta) + \mathscr{L}(\theta) - \hat{\mathscr{L}}(\theta) \right) - \\
&\quad \frac{2}{N} \sup_{\theta \in B_{\delta_{LLN}}} |\log \pi(\theta)| \\
&\geq \inf_{\theta \in B_{\delta_{LLN}} \setminus B_{\delta_2/2}} (\mathscr{L}(\theta_\infty) - \mathscr{L}(\theta)) - 2 \sup_{\theta \in B_{\delta_{LLN}}} \left| \hat{\mathscr{L}}(\theta) - \mathscr{L}(\theta) \right| - \eta \\
&\geq 4\eta - 3\eta = \eta > 0.
\end{aligned}$$

Combining,

$$\inf_{\theta \in R_3} \left( \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta_\infty) - \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta) \right) \geq \min(\eta, \varepsilon) > 0. \tag{S19}$$

Finally, we control the first term in Eq. S17. Since $\hat{\theta}$ maximizes $\hat{\mathscr{L}}(\theta)$, and on the event that $\hat{\theta}$ is in $B_{\delta_{LLN}}$ (which holds under Event 2, since $\left\| \hat{\theta} - \theta_\infty \right\|_2 < \epsilon_U < \delta_{LLN}$),

$$\begin{aligned}
\frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\hat{\theta}) - \frac{1}{N} \sum_{n=1}^{N} \ell(x_n|\theta_\infty) &= \hat{\mathscr{L}}(\hat{\theta}) - \hat{\mathscr{L}}(\theta_\infty) - \frac{1}{N} \log \pi(\hat{\theta}) + \\
&\quad \frac{1}{N} \log \pi(\theta_\infty) \\
&\geq -\frac{2}{N} \sup_{\theta \in B_{\delta_{LLN}}} |\log \pi(\theta)| .
\end{aligned}$$

As argued above in Eq. S18, we can take $N^*$ even larger so that

$$\frac{2}{N} \sup_{\theta \in B_{\delta_{LLN}}} |\log \pi(\theta)| < \frac{1}{2} \min(\eta, \varepsilon).$$

Combining with Eq. S19 and Eq. S17, and multiplying by $N$, gives the desired result with $\epsilon_L = \frac{1}{2}\min(\varepsilon, \eta)$. The argument uses the event in Assumption 1 Item 6, the ULLN event, and Event 2; each holds with probability approaching one, so all three hold jointly with probability approaching one, and we can take $N^*$ large enough that this joint probability is at least $1 - \pi_0$.

□

Note that Lemma S4 requires $\delta_2$ to be sufficiently small (so that $\delta_{LLN} > \delta_2$), which in turn requires $\epsilon_U$ to be small (so that $\delta_{LLN} > 3\epsilon_U$). For any particular $\delta_{LLN} > \delta_2$, we can always apply Lemma S4 with arbitrarily small $\epsilon_U$, at the cost of requiring a large enough $N^*$ so that the probability convergence results apply. Throughout, we will choose $\epsilon_U$ *small enough* that we obtain our desired results.

## D.4 Region $R_3$

The integral over $R_3$ goes to zero by strict optimality of $\theta_\infty$ and the proper behavior of the prior.

**Lemma S5.** *Under Assumption S1, when $\mathfrak{R}$ occurs with $\epsilon_U < \pi(\theta_\infty)/2$,*

$$\|I_3(\phi)\|_2 = O(N^{D_\theta/2}\exp(-N\epsilon_L)) \mathop{\mathbb{E}}_{\pi(\theta)}\left[\|\phi(\theta, \mathcal{X}_s)\|_2\right].$$

*Proof.* We have

$$\begin{aligned}
\|I_3(\phi)\|_2 = N^{D_\theta/2} &\left\| \int_{R_3} \phi(\theta, \mathcal{X}_s) \exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) d\theta \right\|_2 \\
&\leq N^{D_\theta/2} \sup_{\theta \in R_3} \exp\left(N\left(\frac{1}{N}\sum_{n=1}^N \ell(x_n|\theta) - \frac{1}{N}\sum_{n=1}^N \ell(x_n|\hat{\theta})\right)\right) \times \int_{R_3} \|\phi(\theta, \mathcal{X}_s)\|_2 \frac{\pi(\theta)}{\pi(\hat{\theta})} d\theta \\
&\leq N^{D_\theta/2}\exp(-N\epsilon_L)\frac{1}{\pi(\hat{\theta})} \mathop{\mathbb{E}}_{\pi(\theta)}\left[\|\phi(\theta, \mathcal{X}_s)\|_2 \mathbb{I}(\theta \in R_3)\right] \quad \text{(Lemma S4, Event 6)} \\
&\leq N^{D_\theta/2}\exp(-N\epsilon_L)\frac{1}{\pi(\theta_\infty) - \epsilon_U} \mathop{\mathbb{E}}_{\pi(\theta)}\left[\|\phi(\theta, \mathcal{X}_s)\|_2 \mathbb{I}(\theta \in R_3)\right] \quad \text{(Lemma S4, Event 3)} \\
&\leq N^{D_\theta/2}\exp(-N\epsilon_L) 2\pi(\theta_\infty)^{-1} \mathop{\mathbb{E}}_{\pi(\theta)}\left[\|\phi(\theta, \mathcal{X}_s)\|_2\right].
\end{aligned}$$

The factor $N^{D_\theta/2}$ in the first line is the Jacobian of the change of variables Eq. S11, since $I_3(\phi)$ is an integral over $\tau$ (Definition S7). In the last step, we take $\epsilon_U < \pi(\theta_\infty)/2$ in $\mathfrak{R}$ (specifically, Lemma S4 Event 3). □

We note that the form of the bound in Lemma S5 is designed to take advantage of the prior mean assumption Definition 2 Item 2, or its more general version Definition S2 Item 2.

## D.5 Gaussian definitions

For regions $R_1$ and $R_2$, and the combination of the regions, we will make use of the following compact notation for key Gaussian densities.

**Definition S8.** Write the Gaussian density with mean 0 and covariance $\hat{\mathcal{I}}^{-1}$ and its normalizing constant as

$$\hat{\mathcal{G}}(\tau) := \hat{Z}^{-1}\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right) = \mathcal{N}(\tau|0, \hat{\mathcal{I}}^{-1}) \quad \text{where} \quad \hat{Z} := \det\left(2\pi\hat{\mathcal{I}}^{-1}\right)^{1/2}.$$

Additionally, let $\mathcal{G}_\lambda^S(\tau)$ denote the $\mathcal{N}(\tau|0, \lambda^{-1}I_{D_\theta})$ distribution truncated to the set $S$, where $I_{D_\theta}$ denotes the $D_\theta \times D_\theta$ identity matrix. Specifically,

$$\mathcal{G}_\lambda^S(\tau) := \frac{\exp\left(-\frac{\lambda}{2}\|\tau\|_2^2\right)}{\int_S \exp\left(-\frac{\lambda}{2}\|\tau\|_2^2\right) d\tau}.$$

◀

## D.6 Region $R_2$

Region $R_2$ is dealt with by choosing $\delta_2$ sufficiently small that $I_2(\phi)$ is dominated by a particular Gaussian integral. Then we will make that Gaussian integral small by choosing $\delta_1$ sufficiently large.

**Lemma S6.** *Let Assumption S1 hold. Take $\delta_1 \geq \sqrt{24/\lambda_{\mathcal{I}}}$. There exist $\delta_2$ and $\epsilon_U$, defined as functions only of the population quantities $\lambda_{\mathcal{I}}$ and $\mathscr{L}_{(k)}(\theta)$, such that, for any $\mathcal{X}$ such that $\mathfrak{R}$ occurs,*

$$\|I_2(\phi)\|_2 \leq O\left(N^{-3}\right) \mathop{\mathbb{E}}_{\mathcal{G}_{\lambda_{\mathcal{I}}/4}^{R_1 \bigcup R_2}(\tau)} \left[\|\phi(\theta, \mathcal{X}_s)\|_2\right].$$

*Proof.* The main part of the proof consists in showing that there exists a $\delta_2$ sufficiently small so that

$$\|I_2(\phi)\|_2 \leq \int_{R_2} \exp\left(-\frac{\lambda_{\mathcal{I}}}{4}\|\tau\|_2^2\right) \|\phi(\theta, \mathcal{X}_s)\|_2 \, d\tau, \tag{S20}$$

The approach will be to bound the difference between $N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})$ and $\frac{1}{2}\mathscr{L}_{(2)}(\theta_\infty)$ so that the entire expression inside the integral's exponent can be written as $A\tau^2$ for a positive definite matrix $A$.

For the duration of this proof, define the following residual matrix:

$$\begin{aligned}\Lambda(\theta, \mathcal{X}, N) := \frac{1}{2}\left(\hat{\mathscr{L}}_{(2)}(\hat{\theta}) - \mathscr{L}_{(2)}(\theta_\infty)\right) + \\ N^{-1/2}\frac{1}{6}\left(\hat{\mathscr{L}}_{(3)}(\hat{\theta}) - \mathscr{L}_{(3)}(\theta_\infty)\right)\tau + \\ N^{-1/2}\frac{1}{6}\mathscr{L}_{(3)}(\theta_\infty)\tau + N^{-1}\mathscr{T}_{(4)}\left(\hat{\mathscr{L}}, \hat{\theta}, \theta\right)\tau^2.\end{aligned}$$

Under this definition, by Eq. S12,

$$N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta}) = \left(\frac{1}{2}\mathscr{L}_{(2)}(\theta_\infty) + \Lambda(\theta, \mathcal{X}, N)\right)\tau^2.$$

This expansion holds for all $\theta$, but we need to control the residual only in $R_2$. We do so by bounding the entries of the $D_\theta \times D_\theta$ matrix $\Lambda(\theta, \mathcal{X}, N)$. When $\mathfrak{R}$ occurs, we have

$$\begin{aligned}\left\|\hat{\mathscr{L}}_{(2)}(\hat{\theta}) - \mathscr{L}_{(2)}(\theta_\infty)\right\|_2 &\leq \left\|\hat{\mathscr{L}}_{(2)}(\hat{\theta}) - \mathscr{L}_{(2)}(\hat{\theta})\right\|_2 + \left\|\mathscr{L}_{(2)}(\hat{\theta}) - \mathscr{L}_{(2)}(\theta_\infty)\right\|_2 \\ &\leq \sup_{\theta \in R_1 \bigcup R_2}\left\|\hat{\mathscr{L}}_{(2)}(\theta) - \mathscr{L}_{(2)}(\theta)\right\|_2 + \sup_{\theta \in R_1 \bigcup R_2}\left\|\mathscr{L}_{(3)}(\theta)\right\|_2 \left\|\hat{\theta} - \theta_\infty\right\|_2 \\ &\leq \left(1 + \sup_{\theta \in R_1 \bigcup R_2}\left\|\mathscr{L}_{(3)}(\theta)\right\|_2\right)\epsilon_U. \quad \text{(Lemma S4, Events 4 and 2)}\end{aligned}$$

Similarly,

$$\begin{aligned}\left\|\left(\hat{\mathscr{L}}_{(3)}(\hat{\theta})-\mathscr{L}_{(3)}(\theta_\infty)\right)\tau\right\|_2 &\leq \left\|\hat{\mathscr{L}}_{(3)}(\hat{\theta})-\mathscr{L}_{(3)}(\theta_\infty)\right\|_2 \|\tau\|_2 \\ &\leq \left(1+\sup_{\theta\in R_1\bigcup R_2}\left\|\mathscr{L}_{(4)}(\theta)\right\|_2\right)\|\tau\|_2\,\epsilon_U.\end{aligned}$$

Next, applying Cauchy-Schwarz gives

$$\left\|\mathscr{L}_{(3)}(\theta_\infty)\tau\right\|_2 \leq \left\|\mathscr{L}_{(3)}(\theta_\infty)\right\|_2 \|\tau\|_2 .$$

Finally, let us consider the Taylor series residual term. Let $\tilde{\theta}(t)=t(\theta-\hat{\theta})+\hat{\theta}$. Because we can write

$$\begin{aligned}\mathscr{T}_{(4)}\left(\hat{\mathscr{L}},\hat{\theta},\theta\right) &= \frac{1}{6}\int_0^1 (1-t)^3\hat{\mathscr{L}}_{(4)}(\tilde{\theta}(t))dt \\ &= \frac{1}{6}\int_0^1 (1-t)^3\left(\hat{\mathscr{L}}_{(4)}(\tilde{\theta}(t))-\mathscr{L}_{(4)}(\tilde{\theta}(t))\right)dt+\frac{1}{6}\int_0^1 (1-t)^3\mathscr{L}_{(4)}(\tilde{\theta}(t))dt,\end{aligned}$$

we have, when $\mathfrak{R}$ occurs,

$$\begin{aligned}\sup_{\theta\in R_2}\left\|\mathscr{T}_{(4)}\left(\hat{\mathscr{L}},\hat{\theta},\theta\right)\tau^2\right\|_2 \leq& \left\|\mathscr{T}_{(4)}\left(\hat{\mathscr{L}},\hat{\theta},\theta\right)\right\|_2\|\tau\|_2^2 \\ \leq& \frac{1}{6}\left(\sup_{\theta\in R_1\bigcup R_2}\left\|\hat{\mathscr{L}}_{(4)}(\theta)-\mathscr{L}_{(4)}(\theta)\right\|_2+\sup_{\theta\in R_1\bigcup R_2}\left\|\mathscr{L}_{(4)}(\theta)\right\|_2\right)\|\tau\|_2^2 \\ \leq& \frac{1}{6}\left(\epsilon_U+\sup_{\theta\in R_1\bigcup R_2}\left\|\mathscr{L}_{(4)}(\theta)\right\|_2\right)\|\tau\|_2^2. \quad \text{(Lemma S4, Event 4)}\end{aligned}$$

Combining the above results, together with the fact that, in $R_2$, $N^{-1/2}\left\|\tau\right\|_2 \leq \delta_2$, gives

$$\begin{aligned}
\sup_{\theta\in R_2}\left\|\Lambda(\theta,\mathcal{X},N)\right\|_2 \leq & \frac{1}{2}\left\|\hat{\mathscr{L}}_{(2)}(\hat{\theta})-\mathscr{L}_{(2)}(\theta_\infty)\right\|_2 + \\
& N^{-1/2}\frac{1}{6}\left\|\hat{\mathscr{L}}_{(3)}(\hat{\theta})-\mathscr{L}_{(3)}(\theta_\infty)\right\| \sup_{\theta\in R_2}\left\|\tau\right\|_2 + \\
& N^{-1/2}\frac{1}{6}\left\|\mathscr{L}_{(3)}(\theta_\infty)\right\|_2 \sup_{\theta\in R_2}\left\|\tau\right\|_2 + \\
& N^{-1}\sup_{\theta\in R_2}\left\|\mathscr{T}_{(4)}\left(\hat{\mathscr{L}},\hat{\theta},\theta\right)\right\|_2 \left\|\tau\right\|_2^2 \\
\leq & \frac{1}{2}\left(1+\sup_{\theta\in R_1\bigcup R_2}\left\|\mathscr{L}_{(3)}(\theta)\right\|_2\right)\epsilon_U + \\
& \frac{1}{6}\left(1+\sup_{\theta\in R_1\bigcup R_2}\left\|\mathscr{L}_{(4)}(\theta)\right\|_2\right)\delta_2\epsilon_U + \\
& \frac{1}{6}\left\|\mathscr{L}_{(3)}(\theta_\infty)\right\|_2\delta_2 + \frac{1}{6}\left(\epsilon_U+\sup_{\theta\in R_1\bigcup R_2}\left\|\mathscr{L}_{(4)}(\theta)\right\|_2\right)\delta_2^2. \qquad \text{(S21)}
\end{aligned}$$

The preceding bound contains only $\epsilon_U$, $\delta_2$, and the population quantities $\lambda_{\mathcal{I}}$ and $\sup_{\theta\in R_1\bigcup R_2}\left\|\mathscr{L}_{(k)}(\theta)\right\|_2$ for $k\in\{3,4\}$. The latter suprema depend on $\delta_2$ themselves (since $R_1\bigcup R_2$ does), but are finite and decreasing in $\delta_2$.

Consequently, when $\mathfrak{R}$ occurs, there exist an $\epsilon_U$ and $\delta_2$ sufficiently small so that

$$\sup_{\theta\in R_2}\left\|\Lambda(\theta,\mathcal{X},N)\right\|_2 \leq \frac{\lambda_{\mathcal{I}}}{4}, \qquad \text{(S22)}$$

independently of $\mathcal{X}$ (except that $\mathfrak{R}$ occurs). The bound Eq. S22 cannot necessarily be achieved simply by setting $\epsilon_U$ to be small (i.e., by relying on the consistency of $\hat{\theta}$ and the ULLN) because of the terms $\sup_{\theta\in R_1\bigcup R_2}\left\|\mathscr{L}_{(4)}(\theta)\right\|_2\delta_2^2$ and $\left\|\mathscr{L}_{(3)}(\theta_\infty)\right\|_2\delta_2$ in Eq. S21—it is necessary to set $\delta_2$ small as well.

Recall from Definition 1 that $\mathcal{I}=-\mathscr{L}_{(2)}(\theta_\infty)$, and from Definition S4 that $\lambda_{\mathcal{I}}>0$ is the minimum eigenvalue of $\mathcal{I}$. Since $|\Lambda\tau^2|\leq\left\|\Lambda\right\|_{op}\left\|\tau\right\|_2^2\leq\left\|\Lambda\right\|_2\left\|\tau\right\|_2^2$, these facts, combined with Eq. S22, give that, for any $\tau$,

$$\begin{aligned}
\sup_{\theta\in R_2}\left(\frac{1}{2}\mathscr{L}_{(2)}(\theta_\infty)+\Lambda(\theta,\mathcal{X},N)\right)\tau^2 \leq & -\frac{\lambda_{\mathcal{I}}}{2}\left\|\tau\right\|_2^2+\frac{\lambda_{\mathcal{I}}}{4}\left\|\tau\right\|_2^2 \\
= & -\frac{\lambda_{\mathcal{I}}}{4}\left\|\tau\right\|_2^2.
\end{aligned}$$

Plugging into the definition of $I_2(\phi)$ given in Definition S7 gives the desired Eq. S20.

To complete the proof, we use the fact that the normal density is bounded in $R_2$.

$$
\begin{aligned}
\|I_2(\phi)\|_2 &\leq \int_{R_2} \exp\left(-\frac{\lambda_{\mathcal{I}}}{4}\|\tau\|_2^2\right)\|\phi(\theta,\mathcal{X}_s)\|_2\, d\tau \\
&= \int_{R_2} \exp\left(-\frac{\lambda_{\mathcal{I}}}{8}\|\tau\|_2^2 - \frac{\lambda_{\mathcal{I}}}{8}\|\tau\|_2^2\right)\|\phi(\theta,\mathcal{X}_s)\|_2\, d\tau \\
&\leq \sup_{\tau\in R_2}\exp\left(-\frac{\lambda_{\mathcal{I}}}{8}\|\tau\|_2^2\right)\int_{R_2}\exp\left(-\frac{\lambda_{\mathcal{I}}}{8}\|\tau\|_2^2\right)\|\phi(\theta,\mathcal{X}_s)\|_2\, d\tau \\
&\leq \exp\left(-\frac{\lambda_{\mathcal{I}}}{8}\delta_1^2\log N\right)\int_{R_1\bigcup R_2}\exp\left(-\frac{\lambda_{\mathcal{I}}}{8}\|\tau\|_2^2\right)\|\phi(\theta,\mathcal{X}_s)\|_2\, d\tau.
\end{aligned}
$$

In the final line, we have used the fact that $(\log N)^2 \geq \log N$ for all $N \geq 3$. Taking $\delta_1^2 \geq 24/\lambda_{\mathcal{I}}$ and multiplying by the $O(1)$ normalizer for the $\mathcal{G}_{\lambda_{\mathcal{I}}/4}^{R_1\bigcup R_2}(\tau)$ completes the proof. □

By inspecting the proof of Lemma S6, one can see that any polynomial rate of convergence can be achieved by taking $\delta_1$ sufficiently large. Indeed, if the boundary of $R_1$ were of the form $\tau \leq N^k$ for some $k < 1/2$, then $\|I_2(\phi)\|_2$ would decay exponentially and $R_1$ would still shrink to a point in the $\theta$ domain. However, such a rapidly growing boundary between $R_1$ and $R_2$ comes at a cost in the analysis of region $R_1$. In particular, the use of a logarithmically growing boundary is what allows us to give polylog bounds in the subsequent Lemma S10 and in the final result of Theorem 1.

## D.7 Region $R_1$

It is region $R_1$ that will dominate the value of the integral $I(\phi)$. We will first perform a Taylor series expansion of all the terms around their sample versions, and then take probability limits of the resulting expressions.

First, we will show that integrals over $\hat{\mathcal{G}}(\tau)$ in $R_1$ are sufficiently close to the corresponding Gaussian integrals over $\mathbb{R}^{D_\theta}$ for sufficiently large $\delta_1$.

**Lemma S7.** *Let Assumption S1 hold. Let $\tau^k$ represent the $D_\theta$-dimensional array of outer products of $\tau$. When $\mathfrak{R}$ occurs and $\delta_1 \geq \sqrt{24/\lambda_{\mathcal{I}}}$ then, for any $k \geq 0$,*

$$
\int_{R_1}\hat{\mathcal{G}}(\tau)\tau^k d\tau = \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^k\right] + O(N^{-3}).
$$

*Proof.* By definition,

$$
\int_{R_1}\hat{\mathcal{G}}(\tau)\tau^k d\tau - \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^k\right] = \int_{R_2\bigcup R_3}\hat{\mathcal{G}}(\tau)\tau^k d\tau
$$

We will now use a technique similar to that of the proof of Lemma S6 to control the norm of the final integral. Recall that, on $\mathfrak{R}$, the minimum eigenvalue of $\hat{\mathcal{I}}$, $\hat{\lambda}_{\mathcal{I}}$, is at least $\frac{1}{2}\lambda_{\mathcal{I}}$,

and $\left\|\hat{\mathcal{I}}\right\|_{op} < 2\left\|\mathcal{I}\right\|_{op}$ (Lemma S4 Event 5). The latter bound implies that $\hat{Z}^{-1} = O(1)$.

$$
\begin{aligned}
&\left\|\int_{R_2 \bigcup R_3} \hat{\mathcal{G}}(\tau)\tau^k d\tau\right\|_2 \\
&\leq \int_{R_2 \bigcup R_3} \hat{\mathcal{G}}(\tau)\left\|\tau\right\|_2^k d\tau \\
&= \hat{Z}^{-1}\int_{R_2 \bigcup R_3} \exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\left\|\tau\right\|_2^k d\tau \\
&\leq \hat{Z}^{-1}\int_{R_2 \bigcup R_3} \exp\left(-\frac{1}{4}\lambda_{\mathcal{I}}\left\|\tau\right\|_2^2\right)\left\|\tau\right\|_2^k d\tau \\
&\leq \sup_{\tau \in R_2 \bigcup R_3} \exp\left(-\frac{1}{8}\lambda_{\mathcal{I}}\left\|\tau\right\|_2^2\right)\hat{Z}^{-1}\int_{R_2 \bigcup R_3} \exp\left(-\frac{1}{8}\lambda_{\mathcal{I}}\left\|\tau\right\|_2^2\right)\left\|\tau\right\|_2^k d\tau \\
&= \exp\left(-\frac{1}{8}\lambda_{\mathcal{I}}\delta_1^2(\log N)^2\right)\hat{Z}^{-1}\int_{R_2 \bigcup R_3} \exp\left(-\frac{1}{8}\lambda_{\mathcal{I}}\left\|\tau\right\|_2^2\right)\left\|\tau\right\|_2^k d\tau \\
&= O(N^{-3}),
\end{aligned}
$$

where the final line takes $\delta_1^2 \geq \frac{24}{\lambda_{\mathcal{I}}}$, that $\log N < (\log N)^2$ for $N \geq 3$, that $\hat{Z}^{-1} = O(1)$, and that normal moments are finite. □

Next, we series expand the likelihood term.

**Lemma S8.** *Let Assumption S1 hold. For any $\mathcal{X}$ such that $\mathfrak{R}$ occurs, and for $\theta \in R_1$,*

$$
\begin{aligned}
&\exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) = \\
&\qquad \exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right) \times \left(1 + N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 + \tilde{O}\left(N^{-1}\right)\right).
\end{aligned}
$$

*Proof.* Using Eq. S12, write the log likelihood term as

$$
\begin{aligned}
&\exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) = \\
&\qquad \exp\left(\frac{1}{2}\hat{\mathscr{L}}_{(2)}(\hat{\theta})\tau^2\right) \times \\
&\qquad \exp\left(N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 + N^{-1}\mathscr{T}_{(4)}\left(\hat{\mathscr{L}}(\theta), \hat{\theta}, \theta\right)\tau^4\right).
\end{aligned} \tag{S23}
$$

We will expand the second term of Eq. S23 using the expansion of the exponential function around 0, which holds for all $z$:

$$
\begin{aligned}
\exp(z) &= 1 + z + \frac{1}{2}z^2 + \frac{1}{2}\int_0^1 (1-t)^2\exp(tz)dtz^3 \\
&= 1 + z + O(|z|^2).
\end{aligned}
$$

To match Eq. S23, we need to evaluate at

$$z = N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 + N^{-1}\mathscr{T}_{(4)}\left(\hat{\mathscr{L}}(\theta), \hat{\theta}, \theta\right)\tau^4.$$

When $\mathfrak{R}$ occurs (Lemma S4 Event 4), we have

$$\begin{aligned}
\left\| N^{-1}\mathscr{T}_{(4)}\left(\hat{\mathscr{L}}(\theta), \hat{\theta}, \theta\right)\tau^4 \right\|_2 \leq & N^{-1}\sup_{\theta\in R_1}\left\| \mathscr{T}_{(4)}\left(\hat{\mathscr{L}}(\theta), \hat{\theta}, \theta\right)\right\|_2 \sup_{\tau\in R_1}\|\tau\|_2^4 \\
= & \tilde{O}\left(N^{-1}\right) \\
\left\| N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 \right\|_2 \leq & N^{-1/2}\frac{1}{6}\left\|\hat{\mathscr{L}}_{(3)}(\hat{\theta})\right\|_2 \sup_{\tau\in R_1}\left\|\tau^3\right\|_2 \\
= & \tilde{O}\left(N^{-1/2}\right) \Rightarrow \\
|z| = & \tilde{O}\left(N^{-1/2}\right),
\end{aligned}$$

so that

$$\begin{aligned}
\exp\left(N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 + N^{-1}\mathscr{T}_{(4)}\left(\hat{\mathscr{L}}(\theta), \hat{\theta}, \theta\right)\tau^4\right) = \\
1 + N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 + \tilde{O}\left(N^{-1}\right). \qquad \text{(S24)}
\end{aligned}$$

Combining Eqs. S23 and S24 and recognizing the definition of $\hat{\mathcal{I}}$ gives the desired result. □

**Lemma S9.** *Let Assumption S1 hold. When $\mathfrak{R}$ occurs and $\delta_1$ is as given in Lemma S7,*

$$\int_{R_1}\exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right)d\tau = \hat{Z} + \tilde{O}\left(N^{-1}\right).$$

*Proof.* Recall that, by definition,

$$\hat{Z} = \int\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)d\tau.$$

Expanding using Lemma S8 gives

$$
\begin{aligned}
&\left|\int_{R_1} \exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) d\tau - \hat{Z}\right| \\
&= \left|\int_{R_1} \exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\left(1 + N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 + \tilde{O}\left(N^{-1}\right)\right) d\tau - \int \exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right) d\tau\right| \\
&= \left|\int_{R_1} \exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\left(N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3 + \tilde{O}\left(N^{-1}\right)\right) d\tau - \int_{R_2 \bigcup R_3} \exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right) d\tau\right| \\
&\leq \frac{1}{6}N^{-1/2} \sup_{\theta \in R_1} \left\|\hat{\mathscr{L}}_{(3)}(\theta)\right\|_2 \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\mathbb{I}\left(\tau \in R_1\right)\tau^3\right] + \sup_{\tau \in R_2 \bigcup R_3} \exp\left(-\frac{1}{4}\lambda_{\mathcal{I}}\left\|\tau\right\|_2^2\right) + \tilde{O}\left(N^{-1}\right) \\
&= O(N^{-2}) + \exp\left(-\frac{1}{4}\lambda_{\mathcal{I}}\delta_1^2(\log N)^2\right) + \tilde{O}\left(N^{-1}\right) \\
&= \tilde{O}\left(N^{-1}\right),
\end{aligned}
$$

where, in the penultimate line, we use Lemma S7, including the given lower bound on $\delta_1$. □

In Lemma S10 to follow, we integrate the result of Lemma S8 using Lemma S7 to produce our final results for region $R_1$. Recall from Definition S8 that $\hat{\mathcal{G}}(\tau) = \hat{Z}^{-1}\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)$ is the density of the Gaussian distribution on $\tau$ with mean zero and covariance $\hat{\mathcal{I}}^{-1} = (-\hat{\mathscr{L}}_{(2)}(\hat{\theta}))^{-1}$.

Note that the statement of Lemma S10 combines vector-valued quantities like $\phi_{(2)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}[\tau^2]$ with scalar–valued quantities like $\left\|\phi_{(2)}(\hat{\theta}, \mathcal{X}_s)\right\|_2$. This is for notational convenience since the scalar-valued terms will simply become part of the residual, and their component values do not matter. Formally, the equation can be taken to hold componentwise.

**Lemma S10.** *Let Assumption S1 hold. When $\delta_1 \geq \sqrt{24/\lambda_{\mathcal{I}}}$, as in Lemmas S6 and S7, then, for any $\mathcal{X}$ such that $\mathfrak{R}$ occurs,*

$$
\begin{aligned}
I_1(\phi) = \phi(\hat{\theta}, \mathcal{X}_s) &\int_{R_1} \exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) d\tau + \\
&N^{-1}\hat{Z}\left(\frac{1}{2}\phi_{(2)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^2\right] + \frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(1)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^4\right]\right) + \\
&\tilde{O}\left(N^{-3/2}\right)\left(1 + \left\|\phi_{(1)}(\hat{\theta}, \mathcal{X}_s)\right\|_2 + \left\|\phi_{(2)}(\hat{\theta}, \mathcal{X}_s)\right\|_2 + \mathop{\mathbb{E}}_{\mathcal{G}^{R_1}_{\lambda_{\mathcal{I}}/2}(\tau)}\left[\left\|\mathscr{T}_{(3)}\left(\phi(\cdot, \mathcal{X}_s), \hat{\theta}, \theta\right)\right\|_2\right]\right).
\end{aligned}
$$

*Proof.* The result will follow by Taylor expanding $\phi$, combining with Lemma S8, and then integrating using the fact that Gaussian integrals in region $R_1$ are approximately equal to Gaussian integrals over the whole domain, as proved in Lemma S7.

We first Taylor expand $\phi(\theta, \mathcal{X}_s)$ around $\hat{\theta}$ and substitute $\tau = \sqrt{N}(\theta - \hat{\theta})$.

$$
\begin{aligned}
\phi(\theta, \mathcal{X}_s) = \phi(\hat{\theta}, \mathcal{X}_s) &+ N^{-1/2}\phi_{(1)}(\hat{\theta}, \mathcal{X}_s)\tau + \\
&N^{-1}\frac{1}{2}\phi_{(2)}(\hat{\theta}, \mathcal{X}_s)\tau^2 + N^{-3/2}\mathscr{T}_{(3)}\left(\phi(\cdot, \mathcal{X}_s), \hat{\theta}, \theta\right)\tau^3. \qquad \text{(S25)}
\end{aligned}
$$

We will now combine Eq. S25 with Lemma S8.

$$
\begin{aligned}
&\exp\left(N\hat{\mathscr{L}}(\theta)-N\hat{\mathscr{L}}(\hat{\theta})\right)\left(\phi(\theta,\mathcal{X}_s)-\phi(\hat{\theta},\mathcal{X}_s)\right)=\\
&\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\times\\
&\quad\left(N^{-1/2}\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\tau+N^{-1}\frac{1}{2}\phi_{(2)}(\hat{\theta},\mathcal{X}_s)\tau^2+N^{-3/2}\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)\tau^3\right)\times\\
&\quad\left(1+N^{-1/2}\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\tau^3+\tilde{O}\left(N^{-1}\right)\right)=\\
&\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\times\\
&\quad\Bigg(N^{-1/2}\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\tau+\\
&\qquad N^{-1}\left(\frac{1}{2}\phi_{(2)}(\hat{\theta},\mathcal{X}_s)\tau^2+\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\tau^4\right)+\\
&\qquad N^{-3/2}\left(\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)\tau^3+\frac{1}{12}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(2)}(\hat{\theta},\mathcal{X}_s)\tau^5\right)+\\
&\qquad\tilde{O}\left(N^{-3/2}\right)\left(\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\tau+N^{-1/2}\frac{1}{2}\phi_{(2)}(\hat{\theta},\mathcal{X}_s)\tau^2+N^{-1}\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)\tau^3\right)\Bigg)=\\
&\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\times\\
&\quad\Bigg(N^{-1/2}\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\tau+N^{-1}\left(\frac{1}{2}\phi_{(2)}(\hat{\theta},\mathcal{X}_s)\tau^2+\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\tau^4\right)+\\
&\qquad\tilde{O}\left(N^{-3/2}\right)\Bigg(\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\tau+\left(O\left(1\right)\tau^5+O\left(N^{-1/2}\right)\tau^2\right)\phi_{(2)}(\hat{\theta},\mathcal{X}_s)+\\
&\qquad\qquad O\left(1\right)\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)\tau^3\Bigg)\Bigg).
\end{aligned}
\tag{S26}
$$

In the final line of Eq. S26 we absorbed some higher-order quantities into the lower-order $\tilde{O}\left(N^{-3/2}\right)$ term for convenience of expression, and used the fact that, when $\mathfrak{R}$ occurs, for $k=1,2,3$, $\left\|\hat{\mathscr{L}}_{(k)}(\hat{\theta})\right\|_2=O\left(1\right)$ when $\mathfrak{R}$ occurs, since

$$
\left\|\hat{\mathscr{L}}_{(k)}(\hat{\theta})-\mathscr{L}_{(k)}(\hat{\theta})\right\|_2\leq\sup_{\theta\in R_1}\left\|\hat{\mathscr{L}}_{(k)}(\theta)-\mathscr{L}_{(k)}(\theta)\right\|_2\leq\epsilon_U,
$$

so

$$
\begin{aligned}
\left\|\hat{\mathscr{L}}_{(k)}(\hat{\theta})\right\|_2\leq&\left\|\hat{\mathscr{L}}_{(k)}(\hat{\theta})-\mathscr{L}_{(k)}(\hat{\theta})\right\|_2+\left\|\mathscr{L}_{(k)}(\hat{\theta})-\mathscr{L}_{(k)}(\theta_\infty)\right\|_2+\left\|\mathscr{L}_{(k)}(\theta_\infty)\right\|_2\\
\leq&\epsilon_U+\sup_{\theta\in R_1}\left\|\mathscr{L}_{(k)}(\theta)-\mathscr{L}_{(k)}(\theta_\infty)\right\|_2+\left\|\mathscr{L}_{(k)}(\theta_\infty)\right\|_2=O\left(1\right),
\end{aligned}
$$

by continuity of $\mathscr{L}_{(k)}(\theta)$.

Next, we integrate both sides of Eq. S26 using Lemma S7 to replace Gaussian integrals over $R_1$ with Gaussian moments over the whole domain, up to an order $N^{-3}$ which can be absorbed in the $\tilde{O}\left(N^{-3/2}\right)$ error term. (Note that the final term in Eq. S26, $O\left(1\right)\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)\tau^3$, is not an odd function of $\tau$ due to the presence of the Taylor series residuals and so does not vanish in the Gaussian expectation.)

$$\begin{aligned}
I_1(\phi) = &\int_{R_1}\phi(\theta,\mathcal{X}_s)\exp\left(N\hat{\mathscr{L}}(\theta)-N\hat{\mathscr{L}}(\hat{\theta})\right)d\tau \\
= &\,\phi(\hat{\theta},\mathcal{X}_s)\int_{R_1}\exp\left(N\hat{\mathscr{L}}(\theta)-N\hat{\mathscr{L}}(\hat{\theta})\right)d\tau + \\
&N^{-1}\hat{Z}\left(\frac{1}{2}\phi_{(2)}(\hat{\theta},\mathcal{X}_s)\mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^2\right]+\frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^4\right]\right)+ \\
&\tilde{O}\left(N^{-3/2}\right)\left(1+\left\|\phi_{(1)}(\hat{\theta},\mathcal{X}_s)\right\|_2+\left\|\phi_{(2)}(\hat{\theta},\mathcal{X}_s)\right\|_2+\right. \\
&\qquad\left.\int_{R_1}\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)d\tau\right).
\end{aligned}$$

Finally, since $\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)$ depends explicitly on $\tau$, we must control this term in a more convenient form. We will argue as in Lemma S7. Recall that, when $\mathfrak{R}$ occurs, $\hat{\lambda}_{\mathcal{I}}\geq\frac{\lambda_{\mathcal{I}}}{2}$, so that $\hat{Z}=O(1)$ and

$$\begin{aligned}
&\left\|\hat{Z}\int_{R_1}\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)d\tau\right\|_2 \\
&\quad\leq\hat{Z}\int_{R_1}\exp\left(-\frac{1}{2}\hat{\mathcal{I}}\tau^2\right)\left\|\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)\right\|_2 d\tau \\
&\quad\leq\hat{Z}\int_{R_1}\exp\left(-\frac{\lambda_{\mathcal{I}}}{4}\left\|\tau\right\|_2^2\right)\left\|\mathscr{T}_{(3)}\left(\phi(\cdot,\mathcal{X}_s),\hat{\theta},\theta\right)\right\|_2 d\tau.
\end{aligned}$$

The result follows by multiplying by the $O(1)$ normalizer of the $\mathcal{G}^{R_1}_{\lambda_{\mathcal{I}}/2}$ distribution. □

## D.8 Combining the regions

We now combine the results of the previous three sections to evaluate $I(\phi)$. The accuracy of our BCLT expansion of $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}\left[\phi(\theta,\mathcal{X}_s)\right]$ will depend on the following residual function.

**Definition S9.** Let $\lambda_{\mathcal{I}}>0$ denote the minimum eigenvalue of $\mathcal{I}$, and let $\delta>0$ be an arbitrarily small number. For a function $\phi(\theta,\mathcal{X}_s)$ that is third-order BCLT-okay, define the

residual function

$$\begin{aligned}
\mathscr{R}^0(\mathcal{X}_s) &:= 1 + \left\|\phi(\hat{\theta}, \mathcal{X}_s)\right\|_2 + \left\|\phi_{(1)}(\hat{\theta}, \mathcal{X}_s)\right\|_2 + \left\|\phi_{(2)}(\hat{\theta}, \mathcal{X}_s)\right\|_2 \\
\mathscr{R}^1(\mathcal{X}_s) &:= \mathop{\mathbb{E}}_{\mathcal{G}^{R_1}_{\lambda_{\mathcal{I}}/2}(\tau)} \left[\left\| \mathscr{T}_{(3)} \left(\phi(\cdot, \mathcal{X}_s), \hat{\theta}, \theta\right)\right\|_2\right] \\
\mathscr{R}^2(\mathcal{X}_s) &:= \mathop{\mathbb{E}}_{\mathcal{G}^{R_1 \bigcup R_2}_{\lambda_{\mathcal{I}}/4}(\tau)} \left[\|\phi(\theta, \mathcal{X}_s)\|_2\right] \\
\mathscr{R}^3(\mathcal{X}_s) &:= \mathop{\mathbb{E}}_{\pi(\theta)} \left[\|\phi(\theta, \mathcal{X}_s)\|_2\right] \\
\mathscr{R}^\phi(\mathcal{X}_s) &:= \mathscr{R}^0(\mathcal{X}_s) + \mathscr{R}^1(\mathcal{X}_s) + \mathscr{R}^2(\mathcal{X}_s) + \mathscr{R}^3(\mathcal{X}_s).
\end{aligned}$$

◀

**Lemma S11.** *Let Assumption 1 hold, and assume that $\phi$ is third-order BCLT-okay. With $\mathbb{F}$-probability approaching one,*

$$\begin{aligned}
&\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})} [\phi(\theta, \mathcal{X}_s)] - \phi(\hat{\theta}, \mathcal{X}_s) = \\
&\quad N^{-1} \left(\frac{1}{2}\phi_{(2)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)} \left[\tau^2\right] + \frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(1)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)} \left[\tau^4\right]\right) + \\
&\quad \tilde{O}\left(N^{-3/2}\right) \mathscr{R}^\phi(\mathcal{X}_s),
\end{aligned}$$

*where $\mathscr{R}^\phi(\mathcal{X}_s)$ is given in Definition S9.*

*Proof.* First, combine the results of Lemmas S5, S6 and S10, consolidating higher orders into lower orders and expanding the domain of integration in the residual $\mathscr{R}^\phi(\mathcal{X}_s)$.

$$\begin{aligned}
I(\phi) = \phi(\hat{\theta}, \mathcal{X}_s) &\int_{R_1} \exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) d\tau + \\
&N^{-1}\hat{Z} \left(\frac{1}{2}\phi_{(2)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)} \left[\tau^2\right] + \frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(1)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)} \left[\tau^4\right]\right) + \\
&\tilde{O}\left(N^{-3/2}\right) \mathscr{R}^\phi(\mathcal{X}_s). \qquad \text{(S27)}
\end{aligned}$$

For convenience in the final bound of Theorem 1 we have included the term $\left\|\phi(\hat{\theta}, \mathcal{X}_s)\right\|_2$ in the definition of the residual even though it does not occur in Lemma S10.

The final expectation is given by the ratio

$$\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})} [\phi(\theta, \mathcal{X}_s)] = \frac{I(\phi)}{I(1)}.$$

For the remainder of this proof, for compactness of notation, we define

$$\begin{aligned}
\zeta &:= \int_{R_1} \exp\left(N\hat{\mathscr{L}}(\theta) - N\hat{\mathscr{L}}(\hat{\theta})\right) d\tau \\
&= \hat{Z} + \tilde{O}\left(N^{-1}\right) \\
&= \hat{Z}(1 + \tilde{O}\left(N^{-1}\right)),
\end{aligned}$$

which follows from Lemma S9 and the fact that $\hat{Z}$ is bounded away from zero when $\mathfrak{R}$ occurs, since then $\left\|\hat{\mathcal{I}}\right\|_{op} < 2\left\|\mathcal{I}\right\|_{op}$ (Lemma S4 Event 5).

When $\phi(\theta, \mathcal{X}_s) \equiv 1$, we have $I_1(1) = \zeta$ exactly, by the definition of $I_1$ (Definition S7). Further, by Lemmas S6 and S5,

$$\left\|I_2(1)\right\|_2 \leq O\left(N^{-3}\right) \mathop{\mathbb{E}}_{\mathcal{G}_{\lambda_{\mathcal{I}}/4}^{R_1 \bigcup R_2}(\tau)}[1] = O\left(N^{-3}\right) \quad \text{and} \quad \left\|I_3(1)\right\|_2 = O(N^{D_\theta/2}\exp(-N\epsilon_L)) = O\left(N^{-3}\right).$$

Since the regions are disjoint when $\mathfrak{R}$ occurs (Lemma S4 Event 1), $I(1) = I_1(1) + I_2(1) + I_3(1) = \zeta + O(N^{-3})$. Finally, since $\zeta = \hat{Z}(1 + \tilde{O}\left(N^{-1}\right))$ and $\hat{Z}$ is bounded away from zero, $\zeta$ is also bounded away from zero for $N$ sufficiently large, so

$$I(1) = \zeta\left(1 + O\left(N^{-3}\right)\right).$$

We now use the expansion, valid for $|z| < 1$ as $z \to 0$,

$$\frac{1}{1+z} = 1 - z + O(z^2). \tag{S28}$$

Using Eq. S27 for $I(\phi)$ and applying Eq. S28 to $1/I(1)$ and to $1/\zeta$ gives

$$\begin{aligned}
\frac{I(\phi)}{I(1)} &= \frac{I(\phi)}{\zeta(1 + O\left(N^{-3}\right))} \\
&= \frac{I(\phi)}{\zeta} + O\left(N^{-3}\right) I(\phi) \\
&= \phi(\hat{\theta}, \mathcal{X}_s) + N^{-1}\frac{\hat{Z}}{\zeta}\left(\frac{1}{2}\phi_{(2)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^2\right] + \frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(1)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^4\right]\right) + \\
&\quad \tilde{O}\left(N^{-3/2}\right)\mathscr{R}^{\phi}(\mathcal{X}_s) \\
&= \phi(\hat{\theta}, \mathcal{X}_s) + N^{-1}\left(\frac{1}{2}\phi_{(2)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^2\right] + \frac{1}{6}\hat{\mathscr{L}}_{(3)}(\hat{\theta})\phi_{(1)}(\hat{\theta}, \mathcal{X}_s) \mathop{\mathbb{E}}_{\hat{\mathcal{G}}(\tau)}\left[\tau^4\right]\right) + \\
&\quad \tilde{O}\left(N^{-3/2}\right)\mathscr{R}^{\phi}(\mathcal{X}_s).
\end{aligned}$$

Note that we have absorbed terms from the expansion of $I(\phi)$ that are linear in $\left\|\phi_{(1)}(\hat{\theta}, \mathcal{X}_s)\right\|$ and $\left\|\phi_{(2)}(\hat{\theta}, \mathcal{X}_s)\right\|$ into the residual $\mathscr{R}^{\phi}(\mathcal{X}_s)$. To do so we have used the fact that, when $\mathfrak{R}$ occurs, $\hat{\mathscr{L}}_{(3)}(\hat{\theta}) = O(1)$.

Finally, the expansion holds with $\mathbb{F}$-probability approaching one by Lemma S4. □

## D.9 Proof that the BCLT residual is square-summable

All that remains to prove Theorem 1 is to show that the residual is $O_p(1)$ when summed. The form of $\mathscr{R}^{\phi}(\mathcal{X}_s)$ is awkward, but it is designed specifically to commute with suprema over sums, allowing us to apply ULLNs to the remainder of our BCLT, as stated formally in Lemmas S12 and S13 below.

Observe that residual expressions of the form $\sup_{\theta \in R_1 \bigcup R_2}\left\|\phi(\theta, x)\right\|_2$ would not be as

useful, since

$$\frac{1}{N}\sum_{n=1}^{N}\sup_{\theta\in R_1\bigcup R_2}\|\phi(\theta,x_n)\|_2 \geq \sup_{\theta\in R_1\bigcup R_2}\frac{1}{N}\sum_{n=1}^{N}\|\phi(\theta,x_n)\|_2\,,$$

and ULLN bound on the latter does not imply a bound on the former, since the supremum may depend on $x_n$ in a way that causes the sample mean to diverge.

**Lemma S12.** *The Gaussian integrals of Definition S9 commute with sums. For a scalar–valued function $\phi(\theta,\mathcal{X}_s)$ and an index set $\mathcal{S}$,*

$$\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\mathop{\mathbb{E}}_{\mathcal{G}^{R_1\bigcup R_2}_{\lambda_{\mathcal{I}}/4}(\tau)}[\phi(\theta,\mathcal{X}_s)]^2 \leq \sup_{\theta\in R_1\bigcup R_2}\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\phi(\theta,\mathcal{X}_s)^2.$$

*Proof.*

$$\begin{aligned}\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\mathop{\mathbb{E}}_{\mathcal{G}^{R_1\bigcup R_2}_{\lambda_{\mathcal{I}}/4}(\tau)}[\phi(\theta,\mathcal{X}_s)]^2 &\leq \mathop{\mathbb{E}}_{\mathcal{G}^{R_1\bigcup R_2}_{\lambda_{\mathcal{I}}/4}(\tau)}\left[\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\phi(\theta,\mathcal{X}_s)^2\right]\\ &\leq \sup_{\tau\in R_1\bigcup R_2}\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\phi(\hat{\theta}+\tau/\sqrt{N},\mathcal{X}_s)^2\\ &= \sup_{\theta\in R_1\bigcup R_2}\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\phi(\theta,\mathcal{X}_s)^2.\end{aligned}$$

□

**Lemma S13.** *The Taylor series residual of Definition S9 commutes with sums. When $\theta\in R_1\bigcup R_2$, for a k-th order continuously differentiable function $\phi(\theta,\mathcal{X}_s)$ and an index set $\mathcal{S}$,*

$$\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\left\|\mathscr{T}_{(k)}\left(\phi(\cdot,\mathcal{X}_s),\theta,\hat{\theta}\right)\right\|_2^2 \leq \sup_{\theta\in R_1\bigcup R_2}\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\left\|\phi_{(k)}(\theta,\mathcal{X}_s)\right\|_2^2.$$

*Proof.* Recall from Definition S3 that

$$\begin{aligned}&\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\left\|\mathscr{T}_{(k)}\left(\phi(\cdot,\mathcal{X}_s),\theta,\hat{\theta}\right)\right\|_2^2 =\\ &\quad\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\left\|\frac{1}{(k-1)!}\int_0^1(1-t)^{k-1}\phi_{(k)}(\hat{\theta}+t(\theta-\hat{\theta}),\mathcal{X}_s)dt\right\|_2^2 \leq\\ &\quad\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\frac{1}{(k-1)!}\int_0^1\left\|\phi_{(k)}(\hat{\theta}+t(\theta-\hat{\theta}),\mathcal{X}_s)\right\|_2^2 dt =\\ &\quad\frac{1}{(k-1)!}\int_0^1\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\left\|\phi_{(k)}(\hat{\theta}+t(\theta-\hat{\theta}),\mathcal{X}_s)\right\|_2^2 dt \leq\\ &\quad\frac{1}{(k-1)!}\sup_{\theta\in R_1\bigcup R_2}\frac{1}{|\mathcal{S}|}\sum_{s\in\mathcal{S}}\left\|\phi_{(k)}(\theta,\mathcal{X}_s)\right\|_2^2.\end{aligned}$$

The result follows because $k \geq 1$. □

**Lemma S14.** *Under Assumption S1,* $\frac{1}{N}\sum_{n=1}^{N} \mathscr{R}^{\phi}(x_n)^2 = O_p(1)$.

*Proof.* We may take $\delta_2$ to be the smaller of the values required for Definition 2 Item 3 and the region $R_2$ Lemma S6 to hold. Then, by Definition S2 Item 3, for $k = 0, 1, 2, 3$,

$$\sup_{\theta \in R_1 \bigcup R_2} \frac{1}{N}\sum_{n=1}^{N} \left\|\phi_{(k)}(\theta, x_n)\right\|_2^2 = O_p(1),$$

It follows immediately that $\frac{1}{N}\sum_{n=1}^{N} \left\|\phi_{(k)}(\hat{\theta}, x_n)\right\|_2 = O_p(1)$ and $\frac{1}{N}\sum_{n=1}^{N} \left\|\phi_{(k)}(\hat{\theta}, x_n)\right\|_2^2 = O_p(1)$ for $k = 0, 1, 2, 3$. From this and Cauchy-Schwarz it follows that $\frac{1}{N}\sum_{n=1}^{N} \mathscr{R}^{0}(x_n)^2 = O_p(1)$.

Applying Lemma S12 with $\mathcal{S} = \{1, \ldots, N\}$ and $\phi(\theta, \mathcal{X}_s) = \|\phi(\theta, x_n)\|_2$ gives $\frac{1}{N}\sum_{n=1}^{N} \mathscr{R}^{2}(x_n) = O_p(1)$ and $\frac{1}{N}\sum_{n=1}^{N} \mathscr{R}^{2}(x_n)^2 = O_p(1)$.

Similarly, $\frac{1}{N}\sum_{n=1}^{N} \mathscr{R}^{1}(x_n)^2 = O_p(1)$ follows from Lemma S12 and then Lemma S13.

The terms $\frac{1}{N}\sum_{n=1}^{N} \mathscr{R}^{3}(x_n)$ and $\frac{1}{N}\sum_{n=1}^{N} \mathscr{R}^{3}(x_n)^2$ are controlled by assumption in Definition 2 Item 2.

Finally, the fact that $\frac{1}{N}\sum_{n=1}^{N} \mathscr{R}^{\phi}(x_n)^2 = O_p(1)$ follows by Cauchy-Schwarz, since $\mathscr{R}^{\phi}(x_n)$ is the sum of products of terms, the squares of each of which have $O_p(1)$ sum. □

# E Proof of Theorem 2

*Proof.* Throughout the proof, write $\hat{g} := g(\hat{\theta})$ and $\hat{\ell}_n := \ell(x_n|\hat{\theta})$.

Apply Theorem 1 three times, with $\phi_g(\theta) := g(\theta)$, $\phi_\ell(\theta, x_n) := \ell(x_n|\theta)$, and $\phi_{\ell g}(\theta, x_n) := g(\theta)\ell(x_n|\theta)$, all of which are third-order BCLT-okay by Assumption 2. Let the residuals be denoted as $\mathscr{R}^{g}$ (with no $x_n$ dependence), $\mathscr{R}^{\ell}(x_n)$, and $\mathscr{R}^{\ell g}(x_n)$, respectively, and write the leading terms of the expansions of $\phi_g$ and $\phi_\ell$ as

$$L_g := \frac{1}{2} g_{(2)}(\hat{\theta})\hat{\mathcal{I}}^{-1} + \frac{1}{6} g_{(1)}(\hat{\theta})\hat{\mathscr{L}}_{(3)}(\hat{\theta})\hat{\mathcal{M}}$$
$$L_\ell(x_n) := \frac{1}{2} \ell_{(2)}(x_n|\hat{\theta})\hat{\mathcal{I}}^{-1} + \frac{1}{6} \ell_{(1)}(x_n|\hat{\theta})\hat{\mathscr{L}}_{(3)}(\hat{\theta})\hat{\mathcal{M}}.$$

The leading term of Eq. 7 is linear in $\phi$, so by the product rule the leading term of the expansion of $\phi_{\ell g}$ is

$$L_{\ell g}(x_n) := \hat{\ell}_n L_g + \hat{g} L_\ell(x_n) + g_{(1)}(\hat{\theta})\hat{\mathcal{I}}^{-1}\ell_{(1)}(x_n|\hat{\theta}).$$

Under Assumption 1, $\hat{\theta} \xrightarrow[N\to\infty]{prob} \theta_\infty$ and $\hat{\mathcal{I}}^{-1}$ has bounded operator norm with probability approaching one (Lemma S3). These facts, together with the continuity of $g_{(1)}$ and $g_{(2)}$, imply that $\hat{g} = O_p(1)$ and $L_g = O_p(1)$, and, by Cauchy-Schwarz and a ULLN (Assumption 1 Item 5 and Lemma S2), that $\frac{1}{N}\sum_{n=1}^{N} \|L_\ell(x_n)\|_2^2 = O_p(1)$.

The three expansions give

$$\begin{aligned}
\psi_n &= N\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}\left[\phi_{\ell g}(\theta, x_n)\right] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}\left[\phi_g(\theta)\right]\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}\left[\phi_\ell(\theta, x_n)\right]\right) \\
&= N\Big(\hat{g}\hat{\ell}_n + N^{-1}L_{\ell g}(x_n) + \tilde{O}\left(N^{-3/2}\right)\mathscr{R}^{\ell g}(x_n)\Big) - \\
&\quad N\left(\hat{g} + N^{-1}L_g + \tilde{O}\left(N^{-3/2}\right)\mathscr{R}^g\right)\left(\hat{\ell}_n + N^{-1}L_\ell(x_n) + \tilde{O}\left(N^{-3/2}\right)\mathscr{R}^\ell(x_n)\right) \\
&= \left(N\hat{g}\hat{\ell}_n + \hat{\ell}_n L_g + \hat{g}L_\ell(x_n)\right) + g_{(1)}(\hat{\theta})\hat{\mathcal{I}}^{-1}\ell_{(1)}(x_n|\hat{\theta}) + \tilde{O}\left(N^{-1/2}\right)\mathscr{R}^{\ell g}(x_n) - \\
&\quad \left(N\hat{g}\hat{\ell}_n + \hat{\ell}_n L_g + \hat{g}L_\ell(x_n)\right) - N^{-1}L_g L_\ell(x_n) + \tilde{O}\left(N^{-1/2}\right)\left(\hat{g}\mathscr{R}^\ell(x_n) + \hat{\ell}_n\mathscr{R}^g\right) + \\
&\quad \tilde{O}\left(N^{-3/2}\right)\left(L_g\mathscr{R}^\ell(x_n) + L_\ell(x_n)\mathscr{R}^g\right) + \tilde{O}\left(N^{-2}\right)\mathscr{R}^g\mathscr{R}^\ell(x_n) \\
&= g_{(1)}(\hat{\theta})\hat{\mathcal{I}}^{-1}\ell_{(1)}(x_n|\hat{\theta}) + \tilde{O}\left(N^{-1/2}\right)\mathscr{R}^\psi(x_n),
\end{aligned}$$

where the third line uses the definition of $L_{\ell g}(x_n)$, the two terms in large parentheses cancel, and we have combined lower orders into

$$\begin{aligned}
\mathscr{R}^\psi(x_n) := \mathscr{R}^{\ell g}(x_n) + \|\hat{g}\|_2\mathscr{R}^\ell(x_n) + \left|\hat{\ell}_n\right|\mathscr{R}^g + \\
\|L_g\|_2\|L_\ell(x_n)\|_2 + \|L_g\|_2\mathscr{R}^\ell(x_n) + \|L_\ell(x_n)\|_2\mathscr{R}^g + \mathscr{R}^g\mathscr{R}^\ell(x_n).
\end{aligned}$$

We now show that the residual is square summable. By Theorem 1, each of $\frac{1}{N}\sum_{n=1}^N \mathscr{R}^\ell(x_n)^2$, $\frac{1}{N}\sum_{n=1}^N \mathscr{R}^{\ell g}(x_n)^2$, and $\mathscr{R}^g$ are $O_p(1)$. Further, $\frac{1}{N}\sum_{n=1}^N \hat{\ell}_n^2 = O_p(1)$ by Eq. S5 of Lemma S2 with $k = 0$. Combining these with the bounds on $\hat{g}$, $L_g$, and $L_\ell(x_n)$ above, and using $(\sum_{i=1}^7 a_i)^2 \leq 7\sum_{i=1}^7 a_i^2$ to control $\sum_{n=1}^N \mathscr{R}^\psi(x_n)^2$, gives $\frac{1}{N}\sum_{n=1}^N \mathscr{R}^\psi(x_n)^2 = O_p(1)$.

Similarly, $\frac{1}{N}\sum_{n=1}^N \left\|\hat{\mathcal{I}}^{-1}\ell_{(1)}(x_n|\hat{\theta})\right\|_2^2 = O_p(1)$ by a ULLN and the boundedness of the operator norm of $\hat{\mathcal{I}}^{-1}$. By repeated application of Cauchy-Schwarz, we can thus average $\psi_n$ and $\psi_n\psi_n^\intercal$ to get

$$\begin{aligned}
\frac{1}{N}\sum_{n=1}^N \psi_n &= g_{(1)}(\hat{\theta})\hat{\mathcal{I}}^{-1}\frac{1}{N}\sum_{n=1}^N \ell_{(1)}(x_n|\hat{\theta}) + \tilde{O}_p\left(N^{-1/2}\right) \\
&= -N^{-1}g_{(1)}(\hat{\theta})\hat{\mathcal{I}}^{-1}\left(\log\pi\right)_{(1)}(\hat{\theta}) + \tilde{O}_p\left(N^{-1/2}\right) = \tilde{O}_p\left(N^{-1/2}\right) \\
\frac{1}{N}\sum_{n=1}^N \psi_n\psi_n^\intercal &= g_{(1)}(\hat{\theta})\hat{\mathcal{I}}^{-1}\hat{\Sigma}\hat{\mathcal{I}}^{-1}g_{(1)}(\hat{\theta})^\intercal + \tilde{O}_p\left(N^{-1/2}\right) = \hat{V}^{\text{Sand}} + \tilde{O}_p\left(N^{-1/2}\right).
\end{aligned}$$

Here we used the first-order condition $\hat{\mathscr{L}}_{(1)}(\hat{\theta}) = 0$, which gives $\sum_{n=1}^N \ell_{(1)}(x_n|\hat{\theta}) = -\left(\log\pi\right)_{(1)}(\hat{\theta})$. It remains to show that $\hat{V}^{\text{Sand}} \xrightarrow[N\to\infty]{prob} V^g$, which follows from the continuous mapping theorem and the following three facts. First, $g_{(1)}(\hat{\theta}) \xrightarrow[N\to\infty]{prob} g_{(1)}(\theta_\infty)$ by continuity of $g_{(1)}$ and Lemma S3. Second, $\hat{\mathcal{I}} \xrightarrow[N\to\infty]{prob} \mathcal{I}$, as shown in the proof of Lemma S4 Event 5, so that $\hat{\mathcal{I}}^{-1} \xrightarrow[N\to\infty]{prob} \mathcal{I}^{-1}$. Third, each entry of $\ell_{(1)}(x_n|\theta)\ell_{(1)}(x_n|\theta)^\intercal$ is continuous in $\theta$ and bounded in absolute value by $\sup_{\theta\in B_{\delta_{LLN}}}\left\|\ell_{(1)}(x_n|\theta)\right\|_2^2$, which has finite expectation, as shown in the

proof of Lemma S2. Each entry therefore satisfies a ULLN (van der Vaart, 2000, Example 19.8), and, since $\hat{\theta} \xrightarrow[N\to\infty]{prob} \theta_\infty$,

$$\hat{\Sigma} \xrightarrow[N\to\infty]{prob} \mathop{\mathbb{E}}_{\mathbb{F}(x_*)} \left[\ell_{(1)}(x_*|\theta_\infty)\ell_{(1)}(x_*|\theta_\infty)^\intercal\right] = \Sigma,$$

where the final equality uses $\mathop{\mathbb{E}}_{\mathbb{F}(x_*)} \left[\ell_{(1)}(x_*|\theta_\infty)\right] = \mathscr{L}_{(1)}(\theta_\infty) = 0$, which holds because $\theta_\infty$ maximizes $\mathscr{L}(\theta)$ and the interchange of differentiation and expectation is justified (Assumption 1 Item 4).

□

# F Weighting representation

To derive Eq. 10, we first derive a general formula for the second derivative with respect to the data weights $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[g(\theta)]$, and then use the Tower property on the resulting expression. For compactness, let $\ell_n(\theta) = \ell(x_n|\theta)$. Recall that Eq. 9 is a special case of the general formula

$$\begin{aligned}\left.\frac{\partial \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)]}{\partial w_n}\right|_{\mathcal{W}} &= \mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}(a(\theta), \ell_n(\theta)) \\ &= \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)\ell_n(\theta)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)],\end{aligned}$$

which we can apply recursively to itself using the chain rule as follows.

$$
\begin{aligned}
&\left.\frac{\partial^2 \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)]}{\partial w_n \partial w_m}\right|_{\mathcal{W}} \\
&= \mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}(a(\theta)\ell_n(\theta), \ell_m(\theta)) - \\
&\quad \mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}(a(\theta), \ell_m(\theta)) \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)] - \mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}(\ell_n(\theta), \ell_m(\theta)) \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)] \\
&= \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)\ell_n(\theta)\ell_m(\theta)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)\ell_n(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_m(\theta)] - \\
&\quad \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)\ell_m(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)] + \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_m(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)] - \\
&\quad \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)\ell_m(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)] + \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_m(\theta)]\,. \\
&= \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)\ell_n(\theta)\ell_m(\theta)] + 2 \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_m(\theta)] - \\
&\quad \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)\ell_n(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_m(\theta)] - \\
&\quad \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)\ell_m(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)] - \\
&\quad \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)\ell_m(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)] \\
&= \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}\left[(a(\theta) - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[a(\theta)])(\ell_n(\theta) - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_n(\theta)])(\ell_m(\theta) - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\ell_m(\theta)])\right].
\end{aligned}
$$

Now suppose that $\theta = (\gamma, \lambda)$, and that $a(\theta) = \gamma$. By applying the Tower property to the second derivative,

$$
\left.\frac{\partial^2 \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[\gamma]}{\partial w_n^2}\right|_{1_N} = \mathop{\mathbb{E}}_{\mathbb{P}(\gamma|\mathcal{X})}\left[(\gamma - \mathop{\mathbb{E}}_{\mathbb{P}(\gamma|\mathcal{X})}[\gamma]) \mathop{\mathbb{E}}_{\mathbb{P}(\lambda|\mathcal{X},\gamma)}\left[(\ell_n(\gamma,\lambda) - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\ell_n(\gamma,\lambda)])^2\right]\right].
$$

Note that

$$
\begin{aligned}
&\mathop{\mathbb{E}}_{\mathbb{P}(\lambda|\gamma,\mathcal{X})}\left[(\ell_n(\gamma,\lambda) - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\ell_n(\gamma,\lambda)])^2\right] \\
&\quad = \mathop{\mathbb{E}}_{\mathbb{P}(\lambda|\mathcal{X},\gamma)}\left[(\ell_n(\gamma,\lambda) - \mathop{\mathbb{E}}_{\mathbb{P}(\lambda|\mathcal{X},\gamma)}[\ell_n(\gamma,\lambda)] + \mathop{\mathbb{E}}_{\mathbb{P}(\lambda|\mathcal{X},\gamma)}[\ell_n(\gamma,\lambda)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\ell_n(\gamma,\lambda)])^2\right] \\
&\quad = \mathop{\mathrm{Cov}}_{\mathbb{P}(\lambda|\mathcal{X},\gamma)}(\ell_n(\gamma,\lambda)) + \left(\mathop{\mathbb{E}}_{\mathbb{P}(\lambda|\mathcal{X},\gamma)}[\ell_n(\gamma,\lambda)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\ell_n(\gamma,\lambda)]\right)^2.
\end{aligned}
$$

If we assume that the third–order posterior cumulant

$$
\mathop{\mathbb{E}}_{\mathbb{P}(\gamma|\mathcal{X})}\left[(\gamma - \mathop{\mathbb{E}}_{\mathbb{P}(\gamma|\mathcal{X})}[\gamma])\left(\mathop{\mathbb{E}}_{\mathbb{P}(\lambda|\mathcal{X},\gamma)}[\ell_n(\gamma,\lambda)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\ell_n(\gamma,\lambda)]\right)^2\right]
$$

is $\tilde{O}_p\left(N^{-2}\right)$ — as it would be if the corresponding moments of $\mathbb{P}(\gamma|\mathcal{X})$ obeyed an

expansion like Theorem 1 — then Eq. 10 follows.

## F.1 Derivatives of posterior expectations

Consider an exponentially tilted distribution,

$$\mathbb{P}(\theta|t) := \frac{\exp\left(\ell(\theta) + t\phi(\theta)\right)}{\int \exp\left(\ell(\theta') + t\phi(\theta')\right) d\theta'}$$

which we assume is defined for all $t$ in a neighborhood of the zero vector. When we can exchange the order of integration and differentiation,

$$\left.\frac{d \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)]}{dt}\right|_t = \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)\phi(\theta)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[\phi(\theta)] .$$

Applying the previous formula to itself gives

$$\begin{aligned}\left.\frac{d^2 \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)]}{dt^2}\right|_t = & \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)\phi(\theta)\phi(\theta)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)\phi(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[\phi(\theta)] - \\ & \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)] \left( \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[\phi(\theta)\phi(\theta)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[\phi(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[\phi(\theta)] \right) - \\ & \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[\phi(\theta)] \left( \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)\phi(\theta)] - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)] \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[\phi(\theta)] \right) .\end{aligned}$$

If, for some $t_0$, we set

$$\begin{aligned}\bar{g}(\theta) &:= g(\theta) - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t_0)}[g(\theta)] \\ \bar{\phi}(\theta) &:= \phi(\theta) - \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t_0)}[\phi(\theta)] ,\end{aligned}$$

then the above formulas simplify at $t_0$ to

$$\left.\frac{d \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)]}{dt}\right|_{t=t_0} = \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t_0)}\left[\bar{g}(\theta)\bar{\phi}(\theta)\right] \quad \text{and} \quad \left.\frac{d^2 \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t)}[g(\theta)]}{dt^2}\right|_{t=t_0} = \mathop{\mathbb{E}}_{\mathbb{P}(\theta|t_0)}\left[\bar{g}(\theta)\bar{\phi}(\theta)\bar{\phi}(\theta)\right] .$$

Importantly, the validity of the preceding expression for the second derivative requires that the centering is at a fixed $t_0$ — the same $t_0$ at which the derivative is evaluated — and does not depend on $t$.

# G Experiment details and additional experiments

## G.1 Measuring sampling error

Since our experiments use a finite $M$, $\hat{V}^{\text{Bayes}}$ and $\hat{V}^{\text{IJ}}$ are MCMC estimators of $V^{\text{Bayes}}$ and $V^{\text{IJ}}$, respectively. Similarly, sampling with replacement from the rows of $\mathcal{X}$ in Algorithm 2 is equivalent to drawing integer weights $\mathcal{W} \sim \text{Multinomial}(N; N^{-1}, \ldots, N^{-1})$ and computing the weighted posterior expectation $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})}[g(\theta)]$ of Eq. 8. We therefore define the exact bootstrap covariance as the expectation under multinomial resampling,

$$V^{\text{Boot}} := N \mathop{\text{Cov}}_{\mathcal{W}\sim\text{Multinomial}(N;N^{-1},\ldots,N^{-1})} \left( \mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X};\mathcal{W})} [g(\theta)] \right),$$

of which $\hat{V}^{\text{Boot}}$, computed with $B$ bootstrap samples and $M$ MCMC draws for each, is a "doubly Monte-Carlo" estimator. In general, for $B < \infty$, we expect bootstrap sampling variability to remain even as $M \to \infty$. Similarly, for $M < \infty$, we expect MCMC sampling variability to remain even as $B \to \infty$. We will shortly discuss how to estimate these Monte Carlo errors.

Even when the conditions of Theorem 2 are met and the IJ covariance is consistent, we do not expect $V^{\text{IJ}} = V^{\text{Boot}}$, even though $V^{\text{Boot}}$ involves no Monte Carlo error, since we have a finite number $N$ of data points, and the bootstrap and IJ covariance are formally distinct (though each consistent) estimators of $V^g$. The difference $V^{\text{IJ}} - V^{\text{Boot}}$ should decrease as $N$ grows, but, for any finite $N$, the difference can be expected to exceed Monte Carlo error for sufficiently large $M$ and $B$. Similarly, even in the correctly specified case, we do not expect $V^{\text{Bayes}} = V^{\text{IJ}}$, as, again, these are distinct (though consistent) estimators of $g_{(1)}(\theta_\infty)\mathcal{I}^{-1}g_{(1)}(\theta_\infty)^\intercal$.

One might attempt to additionally account for frequentist sampling variability of the original data from the target population in $V^{\text{IJ}}$ or $V^{\text{Boot}}$ as estimators of $V^g$, or of $V^{\text{Bayes}}$ as an estimator of $g_{(1)}(\theta_\infty)\mathcal{I}^{-1}g_{(1)}(\theta_\infty)^\intercal$. However, we expect the quantities $V^{\text{IJ}}$, $V^{\text{Boot}}$, and $V^{\text{Bayes}}$ to be highly correlated under random sampling of the data from the population distribution, and it is necessary to take this correlation into account when computing the variability of, say, the IJ "error" $\hat{V}^{\text{IJ}} - \hat{V}^{\text{Boot}}$. The authors do not see a simple, computationally tractable way to account for this correlation, so we neglect this error in our experiments, and account only for the Monte Carlo error of the preceding paragraph. Since we consider the IJ covariance to succeed when the difference $\hat{V}^{\text{IJ}} - \hat{V}^{\text{Boot}}$ is small relative to random variability, neglecting the frequentist sampling error is conservative.

To estimate the Monte Carlo standard errors, we model the individual entries of the matrices $\hat{V}^{\text{Bayes}}$, $\hat{V}^{\text{IJ}}$, and $\hat{V}^{\text{Boot}}$ as normal distributions with Monte Carlo standard deviations given by the matrices $\Xi^{\text{bayes}}$, $\Xi^{\text{IJ}}$, and $\Xi^{\text{boot}}$, respectively.[1] To estimate $\Xi^{\text{bayes}}$ and $\Xi^{\text{IJ}}$, we block bootstrap the MCMC chains using block sizes that are large relative to estimates of the MCMC mixing times. Note that block bootstrapping here requires only re-arranging

[1]Since both $\hat{V}^{\text{Bayes}}$ and $\hat{V}^{\text{IJ}}$ are based on the same MCMC samples, their Monte Carlo errors are correlated, typically positively. However, to avoid cumbersome exposition of the experimental results, we will report standard errors for $\hat{V}^{\text{Bayes}}$ and $\hat{V}^{\text{IJ}}$ as if they were independent, since accounting for this correlation will tend to reduce the standard errors of the difference $\hat{V}^{\text{Bayes}} - \hat{V}^{\text{IJ}}$, and the difference $\hat{V}^{\text{Bayes}} - \hat{V}^{\text{IJ}}$ tends to be large in our experiments, even relative to the conservative estimate of independent standard errors.

the original MCMC chains, which is fast, and does not require drawing any new MCMC samples. To estimate $\Xi^{\text{boot}}$, we use the fact that each bootstrap sample $\mathscr{G}^b$ from Algorithm 2 contains independent MCMC errors, and so the marginal distribution of the $\mathscr{G}^b$ can be used to estimate the variability both of the bootstrap sampling and the MCMC. We do so using the delta method, as follows. For any two scalar random variables, $t_1$ and $t_2$ drawn according to a joint distribution $p(t_1, t_2)$, we can define $h(a, b, c) := a - bc$ and write

$$\begin{aligned}\operatorname*{Cov}_{p(t_1,t_2)}(t_1, t_2) &= \operatorname*{\mathbb{E}}_{p(t_1,t_2)}[t_1 t_2] - \operatorname*{\mathbb{E}}_{p(t_1,t_2)}[t_1] \operatorname*{\mathbb{E}}_{p(t_1,t_2)}[t_2] \\ &= h\left(\operatorname*{\mathbb{E}}_{p(t_1,t_2)}[t_1 t_2], \operatorname*{\mathbb{E}}_{p(t_1,t_2)}[t_1], \operatorname*{\mathbb{E}}_{p(t_1,t_2)}[t_2]\right).\end{aligned}$$

By applying the delta method to the function $h$, we can estimate the sampling variability of $\operatorname*{Cov}_{p(t_1,t_2)}(t_1, t_2)$ from the covariance of the vector $(t_1 t_2, t_1, t_2)$, which, in turn, can be estimated from draws of $t_1$ and $t_2$. We apply the delta method in this way to each entry of the matrix $\hat{V}^{\text{Boot}}$.

## G.2 Applied regression modeling examples: details

| | Fixed effects only | Random and fixed effects |
|---|---|---|
| Linear regression | 398 | 168 |
| Logistic regression | 192 | 10 |

Total quantities to estimate by model type

| | Cross-covariance estimate | Variance estimate |
|---|---|---|
| Log scale parameter | 191 | 62 |
| Regression parameter | 325 | 190 |

Total quantities to estimate by parameter type

Table S2: Models from Gelman and Hill, 2006

For the ARM experiments of Section 4.2, we needed to make decisions about what was an exchangeable unit. For models with no random effects, we define an exchangeable unit to be a tuple containing a particular response and a regressor. For random effects models, we had to choose an exchangeable unit. When the random effect had many distinct levels, we considered the groups of observations within a level to be a single exchangeable unit. For models whose random effects had few distinct levels, we took the exchangeable unit to either be a more fine-grained partition of the data, or, in some cases treated the observations as fully independent.

Although the posterior covariance is not of primary interest, we compute it to investigate the possibility of differences between the frequentist and posterior covariance estimates. Given that most of the models in the ARM textbook are carefully selected by highly experienced statisticians, one might expect there to be few instances of misspecification. Since the models are additionally mostly low-dimensional, we might expect for the posterior covariance to typically provide a good approximation to the frequentist covariance by the Bernstein–von Mises theorem.

To investigate this intuition, in Figure S1 we reproduce Figure 4 but comparing the bootstrap and Bayesian posterior covariance. The metric plotted is the posterior version of $\Delta_{ij}^{\text{IJ}}$ in Eq. 15:

$$\Delta_{ij}^{\text{Bayes}} := \frac{\hat{V}_{ij}^{\text{Bayes}} - \hat{V}_{ij}^{\text{Boot}}}{\left|\hat{V}_{ij}^{\text{Boot}}\right| + \Xi_{ij}^{\text{boot}}}.$$

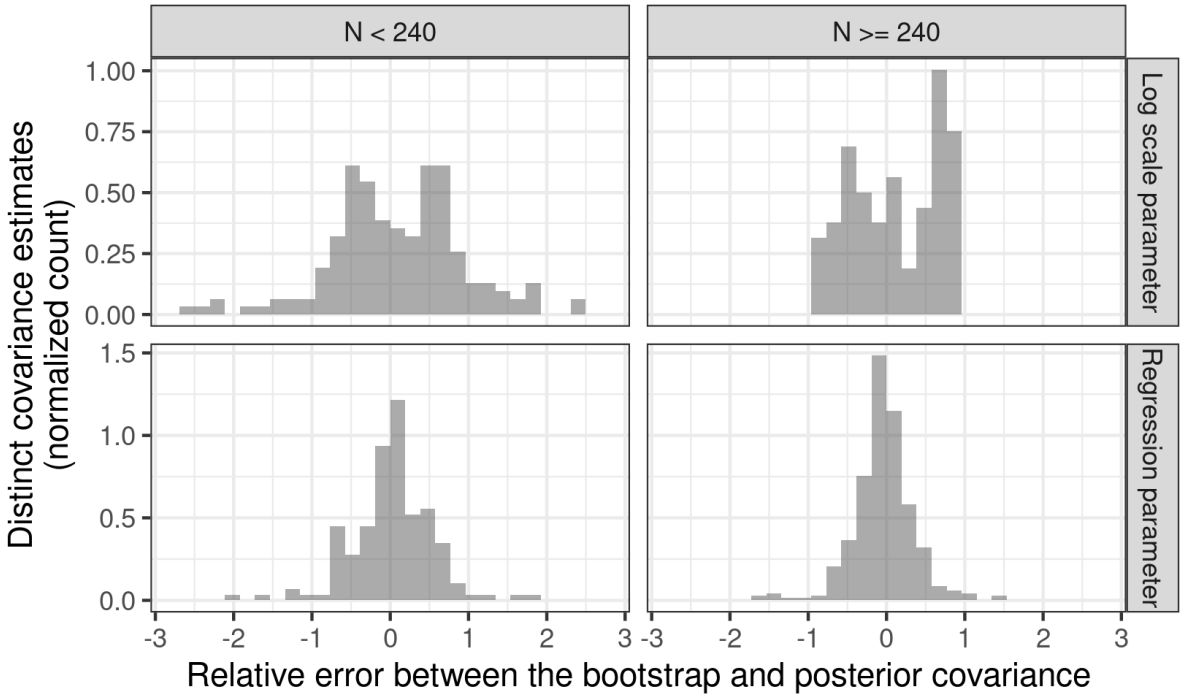



Figure S1: The distribution of the relative errors $\Delta^{\text{Bayes}}$. The "Log scale parameter" panel includes all variances or covariances that involve at least one log scale parameter.

Comparing with Figure 4, we see that the above intuition is mostly correct, and that the frequentist and posterior covariances are largely similar. The notable exception is for log scale parameters in models with relatively large datasets, for which the IJ provides a markedly better approximation to the bootstrap than does the posterior covariance. This suggests the possibility of using the difference between IJ and posterior covariances as a diagnostic for model misspecification, as is sometimes done in the frequentist literature (White, 1982). We leave this possibility as an interesting avenue for future work.

## G.3 Multilevel regression and poststratification

For a more complex model, we consider an application of "multilevel regression and post–stratification," or MrP (Park, Gelman, and Bafumi, 2004). We studied an example from Lopez-Martin, Phillips, and Gelman (2022) in which a subsample of a national survey is used to infer a prediction on the entire population. Since the textbook example involves simulated subsampling, a ground truth is available for comparison.

MrP models can be motivated by the following survey sampling problem. Let $z_*$ denote a vector of regressors, and $y_*$ a binary survey response. Suppose we observe $N_S$ observations

of the form $x_n = (z_n, y_n)$ from a "survey population," and $N_T$ observations consisting only of $z_m$ from a "target population." We assume that $\underset{\mathbb{P}(y_*|z_*)}{\mathbb{E}}[y_*]$ is the same in both the survey and target populations, but that the marginal distributions of $z_*$ may differ, perhaps due to biased survey sampling. Let $\mathbb{P}_T(\cdot)$ denote the target population distribution. We would like to use the observed data to estimate $\underset{\mathbb{P}_T(y_*)}{\mathbb{E}}[y_*] = \underset{\mathbb{P}_T(z_*)}{\mathbb{E}}\left[\underset{\mathbb{P}(y_*|z_*)}{\mathbb{E}}[y_*]\right]$ in the target population. The approach of MrP is to use the survey data to form a regression estimate of $z_* \mapsto \underset{\mathbb{P}(y_*|z_*)}{\mathbb{E}}[y_*]$ (the "multilevel regression" part of MrP) and then to average the regression's predictions using the target population regressors (the "post–stratification" part of MrP).

A typical MrP model assumes that the expectation of the response conditional on the regressors is given by a finitely parameterized regression model. For example, one might specify a logistic regression model for a binary survey response. The regression model gives an explicit expression for $\mathbb{E}[y|z,\theta]$ which, by assumption, holds in both the survey and target populations. Typically, the MrP model also specifies hierarchical, or "multilevel" prior structure on the parameters, possibly including unknown hyperparameters. For compactness of notation we assume that these unknown hyperparameters are also included in $\theta$, even if they are not used directly in the computation of $\mathbb{E}[y|z,\theta]$. For any particular value of $\theta$ we can form the "poststratified" estimator of $\underset{\mathbb{P}_T(y_*)}{\mathbb{E}}[y_*]$, which is our quantity of interest:

$$g(\theta) = \frac{1}{N_T}\sum_{m=1}^{N_T}\mathbb{E}[y_*|z_m,\theta]\,. \tag{S29}$$

The expression for $g(\theta)$ is motivated by the observation that if $\mathbb{E}[y_*|z_m,\theta] = \mathbb{E}[y_*|z_m]$ at our particular $\theta$, then $g(\theta) \approx \underset{\mathbb{P}_T(y_*)}{\mathbb{E}}[y_*]$ by the law of large numbers, since $z_m$ are drawn marginally from the target distribution of regressors. The Bayesian estimator integrates out our uncertainty in the $\theta$, taking $\underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}[g(\theta)]$ as our posterior estimate of the unknown $\underset{\mathbb{P}_T(y_*)}{\mathbb{E}}[y_*]$. See Park, Gelman, and Bafumi (2004) and Lopez-Martin, Phillips, and Gelman (2022) for more details.

To assess our ability to estimate $V^g$ in MrP models, we replicate a semi–simulated analysis from Chapter 1 of the MrP tutorial textbook, Lopez-Martin, Phillips, and Gelman (2022). Although this MrP analysis is based on real data, the survey population consists of a sample taken uniformly at random from a *superpopulation* for which complete data — both regressors and responses — is available. The MrP estimator $\underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}[g(\theta)]$, which is formed using response variables only from the sampled survey population, can then be compared to the known ground truth: namely, the average response in the superpopulation. Lopez-Martin, Phillips, and Gelman (2022) use the superpopulation to assess the accuracy of MrP methods. In our case, we use the superpopulation to estimate the ground truth $V^g$ by measuring the sampling variability of $\underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}[g(\theta)]$ under repeated independent draws of the survey population.

The superpopulation dataset provided by Lopez-Martin, Phillips, and Gelman (2022)

consists of 59,810 pairs of survey responses and regressors, from which we sample 5,000 uniformly at random without replacement to form simulated survey populations. The response $y_*$ is an indicator of whether a respondent supports coverage of abortions in insurance plans, taken from the 2018 Cooperative Congressional Election Study (Schaffner, Ansolabehere, and Luks, 2018). The provided regressors, $z_*$, consist of discretized demographic information measuring age, gender, ethnicity, education, and state. We use the model given by Chapter 1 of Lopez-Martin, Phillips, and Gelman (2022), which consisted of a hierarchical logistic regression with unknown hyperparameter variances. We let our parameter $\theta$ include both logistic regression coefficients and hyperparameter variances; in total, the dimension of the parameter space was $D_\theta = 129$.

The details of our sampling and estimation procedure were as follows. A draw of the survey data $\mathcal{X}$ is produced by selecting 5,000 regressor / response pairs $x_n = (z_n, y_n)$ uniformly at random without replacement from the 59,810 observations in the superpopulation. For each such $\mathcal{X}$, we can compute $\mathbb{P}(\theta|\mathcal{X})$ and so $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[g(\theta)]$, using the 59,810 − 5,000 remaining superpopulation observations as the target population in $g(\theta)$.[2] By simulating multiple draws of $\mathcal{X}$ and computing $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[g(\theta)]$ for each, we can estimate the ground truth $V^g$. In this way, we simulated 101 different survey datasets $\mathcal{X}$. For each dataset, we drew $M = 18{,}000$ MCMC samples from each posterior using the `R` package `brms` (Bürkner, 2018). We then used these samples to compute an MCMC estimate of $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[g(\theta)]$ on each, the variance of which across distinct simulated values of $\mathcal{X}$ provides a ground–truth estimate of $V^g$. On each set of MCMC draws, we computed $\hat{V}^{\mathrm{IJ}}$ and $\mathop{\mathrm{Var}}_{\mathbb{P}(\theta|\mathcal{X})}(g(\theta))$. We randomly chose one of the samples as a basis for 101 bootstrap replicates to compute a bootstrap variance estimate $\hat{V}^{\mathrm{Boot}}$.

As shown in Figure S2, the IJ covariance and bootstrap covariance are nearly identical. In this case, the Bayesian posterior variance also coincides with the frequentist variance. The total computation time for all bootstrap samples was 12 minutes, and the computation for the IJ covariance on the same set of MCMC draws used for the bootstraps was 15 seconds.[3]

As we show in Appendix G.3.1, on this problem the MAP estimate provides a poor estimate of the posterior expectation, likely due to the relatively high dimensionality of $\theta$ and the nonlinearity of the model, particularly through estimation of the random effect variances. Nevertheless, since $N/D_\theta \approx 38.76$, it is plausible that the posterior is concentrated enough in the nuisance parameters to avoid the challenges for IJ in the high–dimensional regime we discussed in Section 3.3. In Appendix G.3.1, we also compare with marginal maximum likelihood estimates provided by the `R` function `glmer` (Bates et al., 2015), noting that `glmer` does not naturally provide estimates of $V^g$ for a complicated quantity depending jointly on the random and fixed effects.

[2]Technically there will be additional variability due to the random selection of the target regressors in Eq. S29, but we are justified in neglecting this variability since the superpopulation is so large.

[3]The computation time for the IJ covariance was dominated by a single call to `brms::posterior_linpred`, which took 10 of the 15 seconds. The function `brms::posterior_linpred` is only computing the log odds for each observation and each MCMC draw, but is doing so with a sequence of nested loops rather than vectorized operations, due to the idiosyncrasies of `rstan` data objects. In a different MCMC environment, this computation could presumably be done much more efficiently.




Figure S2: Covariance estimates for the MrP model. The dots are the estimates computed on the single dataset chosen as the basis for bootstrapping. The boxplots show the spread of the IJ and posterior covariance across simulations, which includes both frequentist and MCMC variability. The horizontal line shows the true sampling variance as estimated by simulation.

### G.3.1 Multilevel regression and poststratification details

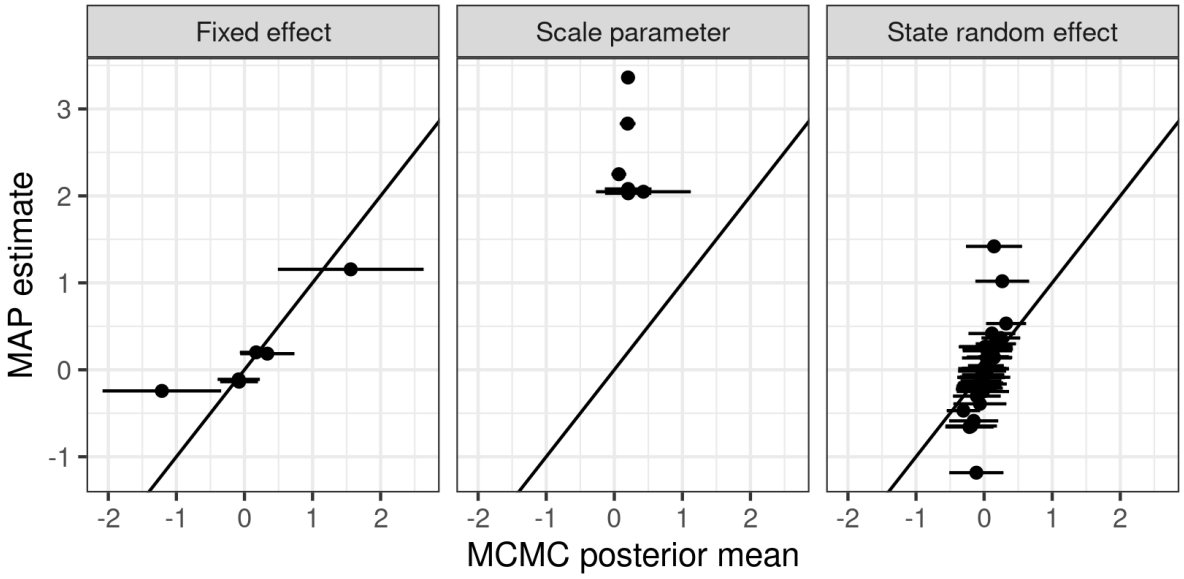



Figure S3: MAP versus MCMC for the MrP model

In Figure S3 we see that the MAP estimator provides poor point estimates in the MrP model, likely due in part to the large number of random effects to fit. We note that the `rstan::optimizing` function in `R` performed poorly, and we had to find the MAP by hand

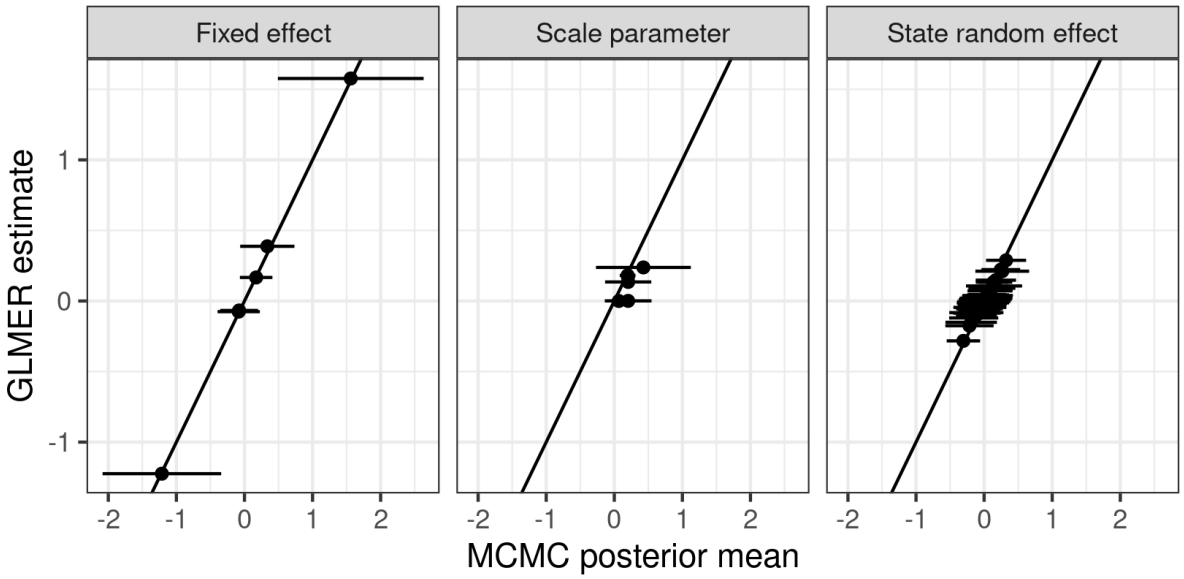



Figure S4: GLMER versus MCMC for the MrP model

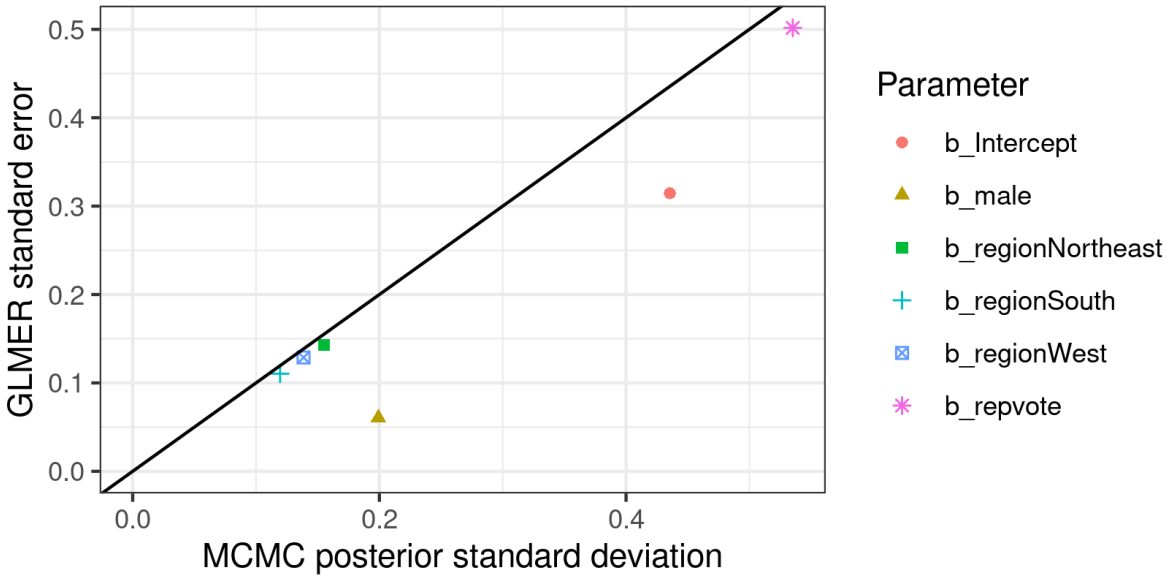



Figure S5: GLMER versus MCMC standard errors for the MrP model

using a careful initialization and small step size in order to avoid divergent estimates.

Although it is not strictly speaking a MAP estimate, the natural optimization–based competitor for MCMC is not the MAP estimator but marginal maximum likelihood estimation, such as the `glmer` function in `R`. As shown in Figure S4, `glmer` provides better point estimates than MAP estimation on the MrP model. However, as shown in Figure S5, the standard error estimates provided by `glmer` can differ considerably from the Bayesian posterior estimates, and most parameters do not have standard errors reported automatically. In principle, one may be able to combine the frequentist sampling estimates of each parameter to estimate the joint sampling distribution between the fixed and random

effects in order to estimate the sampling variance of $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[g(\theta)]$, but doing so would be quite cumbersome given the information exposed by `glmer`. In contrast, the MCMC and IJ covariance estimates are extremely straightforward to compute.

# H Additional closed-form examples

In this section we consider some simple illustrative examples. All the models in this section will be toy problems for which the IJ is not necessary, but they will serve to build intuition for harder, real-life problems.

In each of these examples we will go through the tedious effort to explicitly compute $\psi(x_n)$ and the limit of its sample variance. We stress, however, that we are only doing so for illustrative purposes. In practice, one would simply output $\ell(x_n|\theta)$ from the MCMC algorithm and estimate the posterior covariances using MCMC draws, for which none of the below analytic derivations would be necessary.

### H.0.1 Sample mean

Suppose we observe scalar data, $x_n \in \mathbb{R}$, for $n = 1, \ldots, N$. We assume that $x_n$ and $x_n^2$ obey laws of large numbers, and that $x_n$ itself obeys a central limit theorem, with

$$\begin{aligned}
\hat{\mu} &:= \frac{1}{N}\sum_{n=1}^{N} x_n \xrightarrow[N\to\infty]{prob} \mu \\
\hat{V}_x &:= \frac{1}{N}\sum_{n=1}^{N}(x_n - \hat{\mu})^2 \xrightarrow[N\to\infty]{prob} V_x \\
&\sqrt{N}(\hat{\mu} - \mu) \xrightarrow[N\to\infty]{dist} \mathcal{N}\left(\cdot|0, V_x\right).
\end{aligned}$$

Suppose we use the data to compute a Bayesian posterior using the following model, with fixed prior parameter $\sigma_\theta$, and a fixed variance $\sigma^2$:

$$\begin{aligned}
\theta &\sim \mathcal{N}(\cdot|0, \sigma_\theta^2) \\
x_n|\theta &\overset{iid}{\sim} \mathcal{N}(\cdot|\theta, \sigma^2).
\end{aligned}$$

We will take our parameter of interest to be $\theta$ itself, so $g(\theta) = \theta$. The posterior has a closed form, given by

$$\begin{aligned}
\hat{\theta} &:= \frac{\sigma^{-2}\sum_{n=1}^{N} x_n}{N\sigma^{-2} + \sigma_\theta^{-2}} = \frac{N}{N + \sigma^2/\sigma_\theta^2}\hat{\mu} \\
p(\theta|\mathcal{X}) &= \mathcal{N}\left(\cdot|\hat{\theta}, \frac{\sigma^2}{N + \sigma^2/\sigma_\theta^2}\right).
\end{aligned}$$

Consequently,

$$\begin{aligned}
N \operatorname*{Var}_{\mathbb{P}(\theta|\mathcal{X})}(\theta) &\xrightarrow[N\to\infty]{prob} \sigma^2 \\
\sqrt{N}\left(\operatorname*{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta] - \mu\right) &\xrightarrow[N\to\infty]{dist} \mathcal{N}\left(\cdot|0, V_x\right).
\end{aligned}$$

In this case, if $V_x \neq \sigma^2$—that is, if the model is misspecified— then the limiting variance of the frequentist and Bayesian distributions are different.

How does the IJ perform in this context? In this case, the log likelihood is given up to constants in $\theta$ by

$$\ell(x_n|\theta) = -\frac{1}{2}\sigma^{-2}\left(x_n^2 - 2x_n\theta + \theta^2\right) + Constant.$$

So the influence scores are given by

$$\begin{aligned}\psi_n = & -N\frac{1}{2}\sigma^{-2}\operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})}\left(\theta, -2x_n\theta + \theta^2\right) \\ & = N\sigma^{-2}\operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})}(\theta)\, x_n - N\frac{1}{2}\sigma^{-2}\operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})}\left(\theta, \theta^2\right).\end{aligned}$$

To evaluate the second covariance, use the fact that, by normality,

$$\begin{aligned}\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}\left[(\theta - \hat{\theta})^3\right] = 0 \qquad \text{and} \qquad \theta^2 = (\theta - \hat{\theta})^2 + 2\hat{\theta}\theta - \hat{\theta}^2 \Rightarrow \\ \operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})}\left(\theta, \theta^2\right) = 2\hat{\theta}\operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})}(\theta).\end{aligned}$$

Combining the results, we see that

$$\begin{aligned}\psi_n &= N\sigma^{-2}\operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})}(\theta)\,(x_n - \hat{\theta}) \\ &= \frac{N}{N + \sigma^2/\sigma_\theta^2}\left(x_n - \frac{N}{N + \sigma^2/\sigma_\theta^2}\hat{\mu}\right).\end{aligned}$$

Since $N/(N + \sigma^2/\sigma_\theta^2) \to 1$, the sample variance of $\psi_n$ converges to $V_x$ as desired, despite the fact that the underlying model was misspecified.

### H.0.2 Overdispersed Poisson

Let us now consider another simple non-normal misspecified model, an over-dispersed Poisson distribution.

$$\begin{aligned}\text{Model:} \qquad & \theta \sim \text{Gamma}(a, b) \\ & x_n \overset{iid}{\sim} \text{Poisson}(\theta) \\ \text{Data generating process (reality):} \qquad & \theta = \theta_\infty \\ & \mathop{\mathbb{E}}_{\mathbb{F}(\mathcal{X})}[x_n] = \theta_\infty \\ & \operatorname*{Var}_{\mathbb{F}(\mathcal{X})}(x_n) = \sigma^2 > \theta_\infty.\end{aligned}$$

The posterior has a closed form:

$$\begin{aligned}
\bar{x} &:= \frac{1}{N}\sum_{n=1}^{N} x_n \\
\hat{\sigma}^2 &:= \frac{1}{N}\sum_{n=1}^{N} x_n^2 - \bar{x}^2 \\
\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta] &= \frac{N\bar{x}+a}{N+b} \\
\mathop{\mathrm{Var}}_{\mathbb{P}(\theta|\mathcal{X})}(\theta) &= \frac{N\bar{x}+a}{(N+b)^2} \Rightarrow \\
\mathop{\mathrm{Var}}_{\mathbb{F}(\mathcal{X})}\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta]\right) &= \frac{N\sigma^2}{(N+b)^2}.
\end{aligned}$$

Because $x_n$ is over-dispersed, the posterior variance does not match the frequentist variance $\mathop{\mathrm{Var}}_{\mathbb{F}(\mathcal{X})}\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta]\right)$, even asymptotically.

$$\begin{aligned}
N\mathop{\mathrm{Var}}_{\mathbb{F}(\mathcal{X})}\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\theta]\right) &\xrightarrow[N\to\infty]{prob} \sigma^2 \\
N\mathop{\mathrm{Var}}_{\mathbb{P}(\theta|\mathcal{X})}(\theta) &\xrightarrow[N\to\infty]{prob} \theta_\infty < \sigma^2.
\end{aligned}$$

The model and MAP are given by

$$\begin{aligned}
\log \pi(\theta) &= (a-1)\log\theta - b\theta + C \\
\hat{\mathscr{L}}(\theta) &= \frac{1}{N}\sum_{n=1}^{N}(x_n \log\theta - \theta) + \frac{1}{N}\log\pi(\theta) + C \Rightarrow \\
\hat{\theta} &= \frac{N\bar{x}+a-1}{N+b} \\
\ell_{(1)}(x_n|\theta) &= \frac{x_n}{\theta} - 1 \\
\hat{\mathscr{L}}_{(2)}(\theta) &= -\frac{\bar{x}+(a-1)/N}{\theta^2}.
\end{aligned}$$

Here, the MAP and posterior expectation differ by $1/(N+b)$, though, as expected, the difference is $o_p(N^{-1/2})$. We can now compute the sandwich covariance estimator:

$$\begin{aligned}
\hat{\Sigma} &= \frac{1}{N}\sum_{n=1}^{N}\left(\frac{x_n}{\hat{\theta}} - 1\right)^2 = \frac{\hat{\sigma}^2 + (\bar{x}-\hat{\theta})^2}{\hat{\theta}^2} \\
\hat{\mathcal{I}} &= -\hat{\mathscr{L}}_{(2)}(\hat{\theta}) = \frac{N+b}{N\hat{\theta}} \\
\hat{V}^g &= \hat{\mathcal{I}}^{-1}\hat{\Sigma}\hat{\mathcal{I}}^{-1} = \frac{N^2}{(N+b)^2}\left(\hat{\sigma}^2 + (\bar{x}-\hat{\theta})^2\right).
\end{aligned}$$

As expected, the limits of $N \underset{\mathbb{F}(\mathcal{X})}{\text{Var}}\left(\underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}[\theta]\right)$ and $\hat{V}^g$ are the same.

The influence function is given by

$$\begin{aligned}\psi(x_n) &= N \underset{\mathbb{P}(\theta|\mathcal{X})}{\text{Cov}}(\theta, x_n \log\theta - \theta) \\ &= N x_n \underset{\mathbb{P}(\theta|\mathcal{X})}{\text{Cov}}(\theta, \log\theta) - N \underset{\mathbb{P}(\theta|\mathcal{X})}{\text{Var}}(\theta) \\ &= \frac{N}{N+b} x_n + C.\end{aligned}$$

where $C$ is the same for every $x_n$, and where we have used the posterior covariance $\underset{\mathbb{P}(\theta|\mathcal{X})}{\text{Cov}}(\theta, \log\theta)$ for the Gamma distribution, which can be found by taking the mixed second-order partial derivative of the log partition function. It follows that

$$V^{\text{IJ}} = \frac{1}{N}\sum_{n=1}^{N}\psi(x_n)^2 - \bar{\psi}^2 = \frac{N^2}{(N+b)^2}\hat{\sigma}^2,$$

which is also a consistent estimator of $\sigma^2 = V^g$. Since $\bar{x} - \hat{\theta} = O_p(N^{-1})$, $V^{\text{IJ}}$ and $\hat{V}^g$ differ only by a term of order $O_p(N^{-2})$.

### H.0.3 Multivariate regression

We now consider multivariate regression with unknown variance. This will provide a simple multivariate example, draw a connection with robust frequentist estimators, and allow us to discuss a point about conditioning on part of the exchangeable unit.

The model and data generating process are

Model:
$$\begin{aligned}\sigma^{-2} &\sim \text{Gamma}(a, b) \\ \beta|\sigma &\sim \mathcal{N}\left(\cdot \,\middle|\, 0, \frac{\sigma^2}{\sigma_\beta^2} I_D\right) \\ y_n &\overset{iid}{\sim} \mathcal{N}(\cdot|\beta^T z_n, \sigma^2)\end{aligned}$$

Data generating process (reality): The moments of $y_n z_n$, $z_n z_n^T$, $y_n^2 z_n z_n^T$, and $\|z_n\|_2^4$ are finite and obey a LLN, with invertible $\underset{\mathbb{F}(\mathcal{X})}{\mathbb{E}}\left[z_n z_n^T\right]$.

Here, $\theta = (\beta, \sigma)$ and $x_n = (y_n, z_n)$. We will consider $g(\theta) = \beta$. The full posterior has a

closed form, but we will only require the conditional posterior $p(\beta|x,\sigma)$.

$$\begin{aligned}
\hat{S}_z &:= \frac{1}{N}\sum_{n=1}^{N} z_n z_n^T + \frac{\sigma_\beta^2}{N} I_D \\
\hat{\beta} &:= \hat{S}_z^{-1}\left(\frac{1}{N}\sum_{n=1}^{N} z_n y_n\right) \\
p(\beta|x,\sigma) &= \mathcal{N}\left(\beta|\hat{\beta}, \frac{\sigma^2}{N}\hat{S}_z^{-1}\right) \\
\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\beta] &= \hat{\beta} \\
\mathop{\mathrm{Cov}}_{p(\beta|x,\sigma)}(\beta) = \frac{\sigma^2}{N}\hat{S}_z^{-1} \quad &\text{and} \quad \mathop{\mathrm{Cov}}_{\mathbb{P}(\theta|\mathcal{X})}(\beta) = \frac{\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\sigma^2]}{N}\hat{S}_z^{-1}.
\end{aligned}$$

In order to consider the frequentist variance of $\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\beta]$, we have to specify what is random. Since $x_n = (z_n, y_n)$, the analysis of this paper appears to require that both $z_n$ and $y_n$ are random. However, the generative process for the regressors $z_n$ are not part of the Bayesian model given above. Indeed, the generating process for the regressors is hardly ever specified in Bayesian regression because the data generating process for $z_n$ is ancillary to the parameters of interest, $\theta$. However, precisely because $p(z_n)$ is ancillary for $\theta$, one could resolve this apparent inconsistency by imagining including any data generating process $p(z_n)$ to the Bayesian model, and observing that the Bayesian posterior is unchanged no matter what $p(z_n)$ is.

Asymptotically, the difference between random $z_n$ and fixed $z_n$ does not matter since sample moments over $z_n$ converge to their population counterparts. However, in finite sample, the difference can be important, and it is easier to bootstrap pairs $(z_n, y_n)$ than keep $z_n$ fixed and resample only $y_n$. (We will not consider methods such as the residual bootstrap here.) Unless otherwise stated, we will thus take the view in the present paper that the full exchangeable pair is random.

Because of the term $\hat{S}_z^{-1}$ and the randomness of $z_n$, the variance $N\mathop{\mathrm{Var}}_{\mathbb{F}(\mathcal{X})}\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\beta]\right)$ does not have a closed form. However, by a CLT applied to $\sqrt{N}\frac{1}{N}\sum_{n=1}^{N} z_n(y_n - \beta_\infty^T z_n)$, where $\beta_\infty := \mathop{\mathbb{E}}_{\mathbb{F}(\mathcal{X})}\left[z_n z_n^T\right]^{-1}\mathop{\mathbb{E}}_{\mathbb{F}(\mathcal{X})}[z_n y_n]$, and Slutsky's theorem, the variance is consistently estimated by

$$\begin{aligned}
\hat{\epsilon}_n &:= y_n - \hat{\beta}^T z_n \\
N\mathop{\mathrm{Var}}_{\mathbb{F}(\mathcal{X})}\left(\mathop{\mathbb{E}}_{\mathbb{P}(\theta|\mathcal{X})}[\beta]\right) &\approx \hat{S}_z^{-1}\left(\frac{1}{N}\sum_{n=1}^{N}\hat{\epsilon}_n^2 z_n z_n^T\right)\hat{S}_z^{-1}. \qquad \text{(S30)}
\end{aligned}$$

If the model is correctly specified, and $\sigma = \sigma_0$, then

$$\frac{1}{N}\sum_{n=1}^{N} z_n z_n^T \hat{\epsilon}_n^2 \xrightarrow[N\to\infty]{prob} \sigma_0^2 \mathop{\mathbb{E}}_{\mathbb{F}(\mathcal{X})}\left[z_n z_n^T\right],$$

and $N \underset{\mathbb{F}(\mathcal{X})}{\mathrm{Var}}\left(\underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}[\beta]\right)$ coincides asymptotically with $N \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathrm{Cov}}(\beta)$. However, if the variances of the residuals are correlated with $z_n$, or the mean of $y_n$ is not linear in $z_n$, then the frequentist and Bayesian covariance will not match, even asymptotically. The covariance estimate in Eq. S30 is known in the frequentist regression literature as a "robust" covariance estimate, and it is sometimes recommended when there is concern about model misspecification (e.g. Davidson and MacKinnon, 2004, Section 5.5). In the present context, it is simply the regression version of the plug-in sandwich covariance estimate $\hat{V}^g$ of $V^g$.

The influence function in this case is given by

$$\begin{aligned}
\psi(x_n) &= N \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathrm{Cov}}\left(\beta, -\log\sigma - \frac{1}{2\sigma^2}\left(y_n^2 - 2\beta^T z_n y_n + \beta^T z_n z_n^T \beta\right)\right) \\
&= N \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathrm{Cov}}\left(\beta, \sigma^{-2}\beta\right) z_n y_n - N \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathrm{Cov}}\left(\beta, \log\sigma\right) - \frac{N}{2} \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathrm{Cov}}\left(\beta, \sigma^{-2}\right) y_n^2 - \\
&\quad \frac{N}{2} \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathrm{Cov}}\left(\beta, \sigma^{-2}\beta^T z_n z_n^T \beta\right).
\end{aligned}$$

First,

$$\begin{aligned}
\underset{\mathbb{P}(\theta|\mathcal{X})}{\mathrm{Cov}}\left(\beta, \sigma^{-2}\beta\right) &= \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[\left(\beta - \hat{\beta}\right)\left(\sigma^{-2}\beta - \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[\sigma^{-2}\beta\right]\right)\right] \\
&= \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[\left(\beta - \hat{\beta}\right)\sigma^{-2}\beta\right] \\
&= \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[\sigma^{-2} \underset{p(\beta|x,\sigma)}{\mathbb{E}}\left[\left(\beta - \hat{\beta}\right)\beta\right]\right] \\
&= \frac{1}{N} \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[\sigma^{-2}\sigma^2 \hat{S}_z^{-1}\right] \\
&= \frac{1}{N}\hat{S}_z^{-1}.
\end{aligned}$$

Now, the conditional posterior mean $\underset{p(\beta|x,\sigma)}{\mathbb{E}}[\beta] = \hat{\beta}$ does not depend on $\sigma$, so, for any function $f(\sigma)$ with finite posterior variance (in particular, for $f(\sigma) = \sigma^{-2}$ and $f(\sigma) = \log\sigma$),

$$\begin{aligned}
\underset{\mathbb{P}(\theta|\mathcal{X})}{\mathrm{Cov}}\left(\beta, f(\sigma)\right) &= \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[\left(\beta - \hat{\beta}\right)\left(f(\sigma) - \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[f(\sigma)\right]\right)\right] \\
&= \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[\underset{p(\beta|x,\sigma)}{\mathbb{E}}\left[\beta - \hat{\beta}\right]\left(f(\sigma) - \underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[f(\sigma)\right]\right)\right] \\
&= 0.
\end{aligned}$$

Similarly, the third-order array of third-order covariances is identically zero:

$$\underset{\mathbb{P}(\theta|\mathcal{X})}{\mathbb{E}}\left[\sigma^{-2}(\beta - \hat{\beta})(\beta - \hat{\beta})(\beta - \hat{\beta})\right] = 0.$$

To use this in $\psi(x_n)$, observe that

$$\beta^T z_n z_n^T \beta = \left(\beta - \hat{\beta}\right)^T z_n z_n^T \left(\beta - \hat{\beta}\right) + 2\hat{\beta}^T z_n z_n^T \beta - \hat{\beta}^T z_n z_n^T \hat{\beta},$$

so that

$$\begin{aligned}
\frac{N}{2} \operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})} \left(\beta, \sigma^{-2}\beta^T z_n z_n^T \beta\right) &= N \operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})} \left(\beta, \sigma^{-2}\hat{\beta}^T z_n z_n^T \beta\right) \\
&= N \operatorname*{Cov}_{\mathbb{P}(\theta|\mathcal{X})} \left(\beta, \sigma^{-2}\beta\right) z_n z_n^T \hat{\beta} \\
&= \hat{S}_z^{-1} z_n z_n^T \hat{\beta}.
\end{aligned}$$

Combining,

$$\psi(x_n) = \hat{S}_z^{-1} z_n \left(y_n - z_n^T \hat{\beta}\right).$$

We thus see that

$$\begin{aligned}
\bar{\psi} &= \hat{S}_z^{-1} \frac{1}{N} \sum_{n=1}^{N} z_n \hat{\epsilon}_n = \frac{\sigma_\beta^2}{N} \hat{S}_z^{-1} \hat{\beta} = O_p(N^{-1}) \\
V^{\mathrm{IJ}} &= \frac{1}{N} \sum_{n=1}^{N} \psi(x_n)\psi(x_n)^T - \bar{\psi}\bar{\psi}^T = \hat{S}_z^{-1} \left(\frac{1}{N} \sum_{n=1}^{N} \hat{\epsilon}_n^2 z_n z_n^T\right) \hat{S}_z^{-1} + O_p(N^{-2}),
\end{aligned}$$

where the expression for $\bar{\psi}$ follows from the definition of $\hat{\beta}$, which gives $\frac{1}{N}\sum_{n=1}^{N} z_n y_n = \hat{S}_z \hat{\beta}$. So $V^{\mathrm{IJ}}$ matches the frequentist covariance estimate given in Eq. S30 up to a term of order $O_p(N^{-2})$.

# References


Bates, D., M. Mächler, B. Bolker, and S. Walker (2015). "Fitting Linear Mixed-Effects Models Using lme4". In: *Journal of Statistical Software* 67.1, pp. 1–48. DOI: 10.18637/jss.v067.i01.

Bürkner, P. (2018). "Advanced Bayesian Multilevel Modeling with the R Package brms". In: *The R Journal* 10.1, pp. 395–411. DOI: 10.32614/RJ-2018-017.

Davidson, R. and J. MacKinnon (2004). *Econometric theory and methods*. Vol. 5. Oxford University Press New York.

Dudley, R. (2018a). *Real analysis and probability*. Chapman and Hall/CRC.

— (2018b). *Real analysis and probability*. Chapman and Hall/CRC.

Gelman, A. and J. Hill (2006). *Data analysis using regression and multilevel/hierarchical models*. Cambridge university press.

Huber, P. J. (2011). "Robust statistics". In: *International Encyclopedia of Statistical Science*. Springer, pp. 1248–1251.

Keener, R. (2010). *Theoretical statistics: Topics for a core course*. Springer Science & Business Media.

Lopez-Martin, J., J. Phillips, and A. Gelman (2022). *Multilevel Regression and Poststratification Case Studies*. URL: https://bookdown.org/jl5522/MRP-case-studies/ (visited on 03/05/2024).

Park, D., A. Gelman, and J. Bafumi (2004). "Bayesian multilevel estimation with poststratification: State-level estimates from national polls". In: *Political Analysis* 12.4, pp. 375–385.

Schaffner, B., S. Ansolabehere, and S. Luks (2018). *CCES 2018 Harvard Dataverse*. URL: https://doi.org/10.7910/DVN/ZSBZ7K.

van der Vaart, A. (2000). *Asymptotic statistics*. Vol. 3. Cambridge university press.

van der Vaart, A. and J. Wellner (2013). *Weak convergence and empirical processes with applications to statistics*. Springer Science & Business Media.

White, H. (1982). "Maximum likelihood estimation of misspecified models". In: *Econometrica: Journal of the econometric society*, pp. 1–25.